



VMTN6718U

# Hackathon Training: Open Source Projects for Enterprise Cloud Native

Henry Zhang, Chief Architect, R&D, VMware China

Ben Corrie, Senior Staff Engineer

Patrick Daigle, Sr. Tech Marketing Architect, VIC

Steven Zou, Staff Engineer, Project Harbor

#Vmworld #VMTN6718U



vmworld<sup>®</sup> 2017

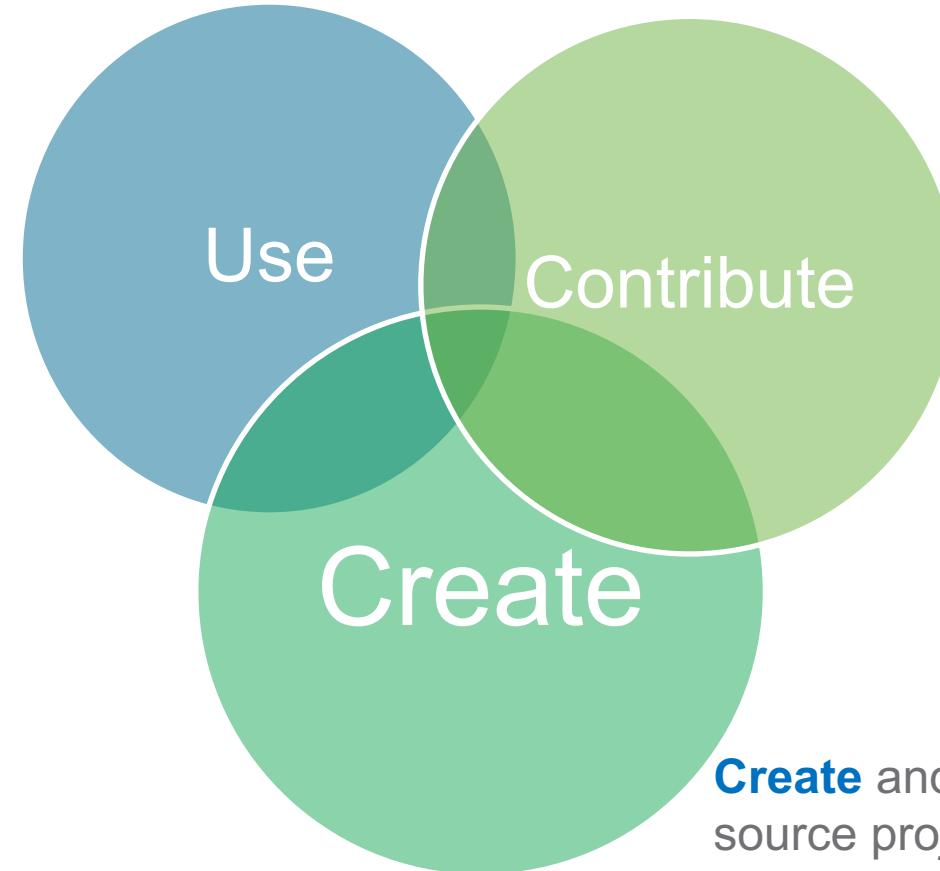
# Disclaimer

- This presentation may contain product features that are currently under development.
- This overview of new technology represents no commitment from VMware to deliver these features in any generally available product.
- Features are subject to change, and must not be included in contracts, purchase orders, or sales agreements of any kind.
- Technical feasibility and market demand will affect final delivery.
- Pricing and packaging for any new technologies or features discussed or presented have not been determined.

# VMware's Participation in Open Source Community

**Use** open source code within the enterprise:

- For implementation of industry standards
- To improve software usability
- To accelerate our software development processes



**Contribute** to upstream projects to add features or fix bugs:

- Linux kernel drivers for our virtual hardware, but also work on key kernel infrastructure
- Open Stack improvement for network, storage, stability
- Kubernetes integration with vSphere as compute and storage provider
- Many smaller contributions and bug fixes to a wide variety of projects

**Create** and release a new open source project and build new community to support it

- Variety of project – from utilities and glue code to significant, standard-setting projects

# Leadership in Open Source Cloud Native Projects



vSphere Integrated  
Containers Engine



HARBOR™

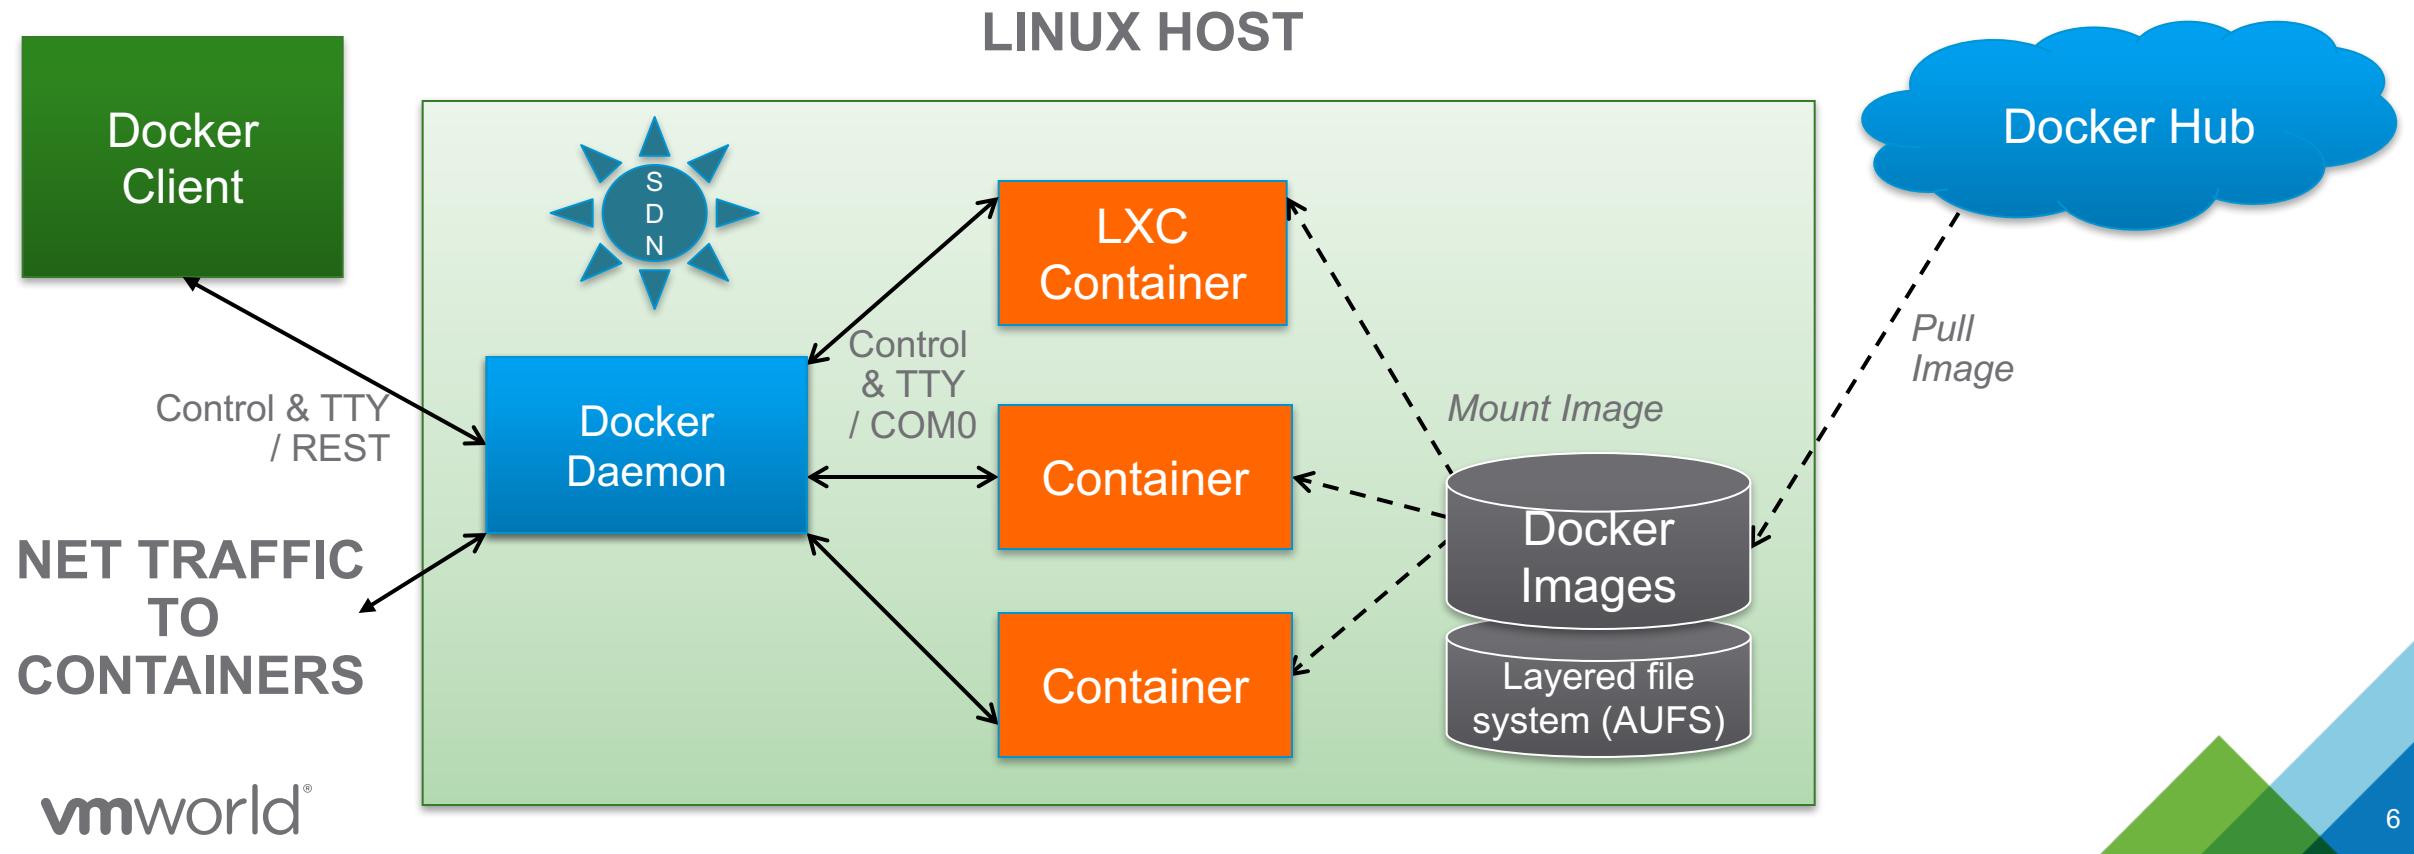


ADMIRAL™

# vSphere Integrated Containers

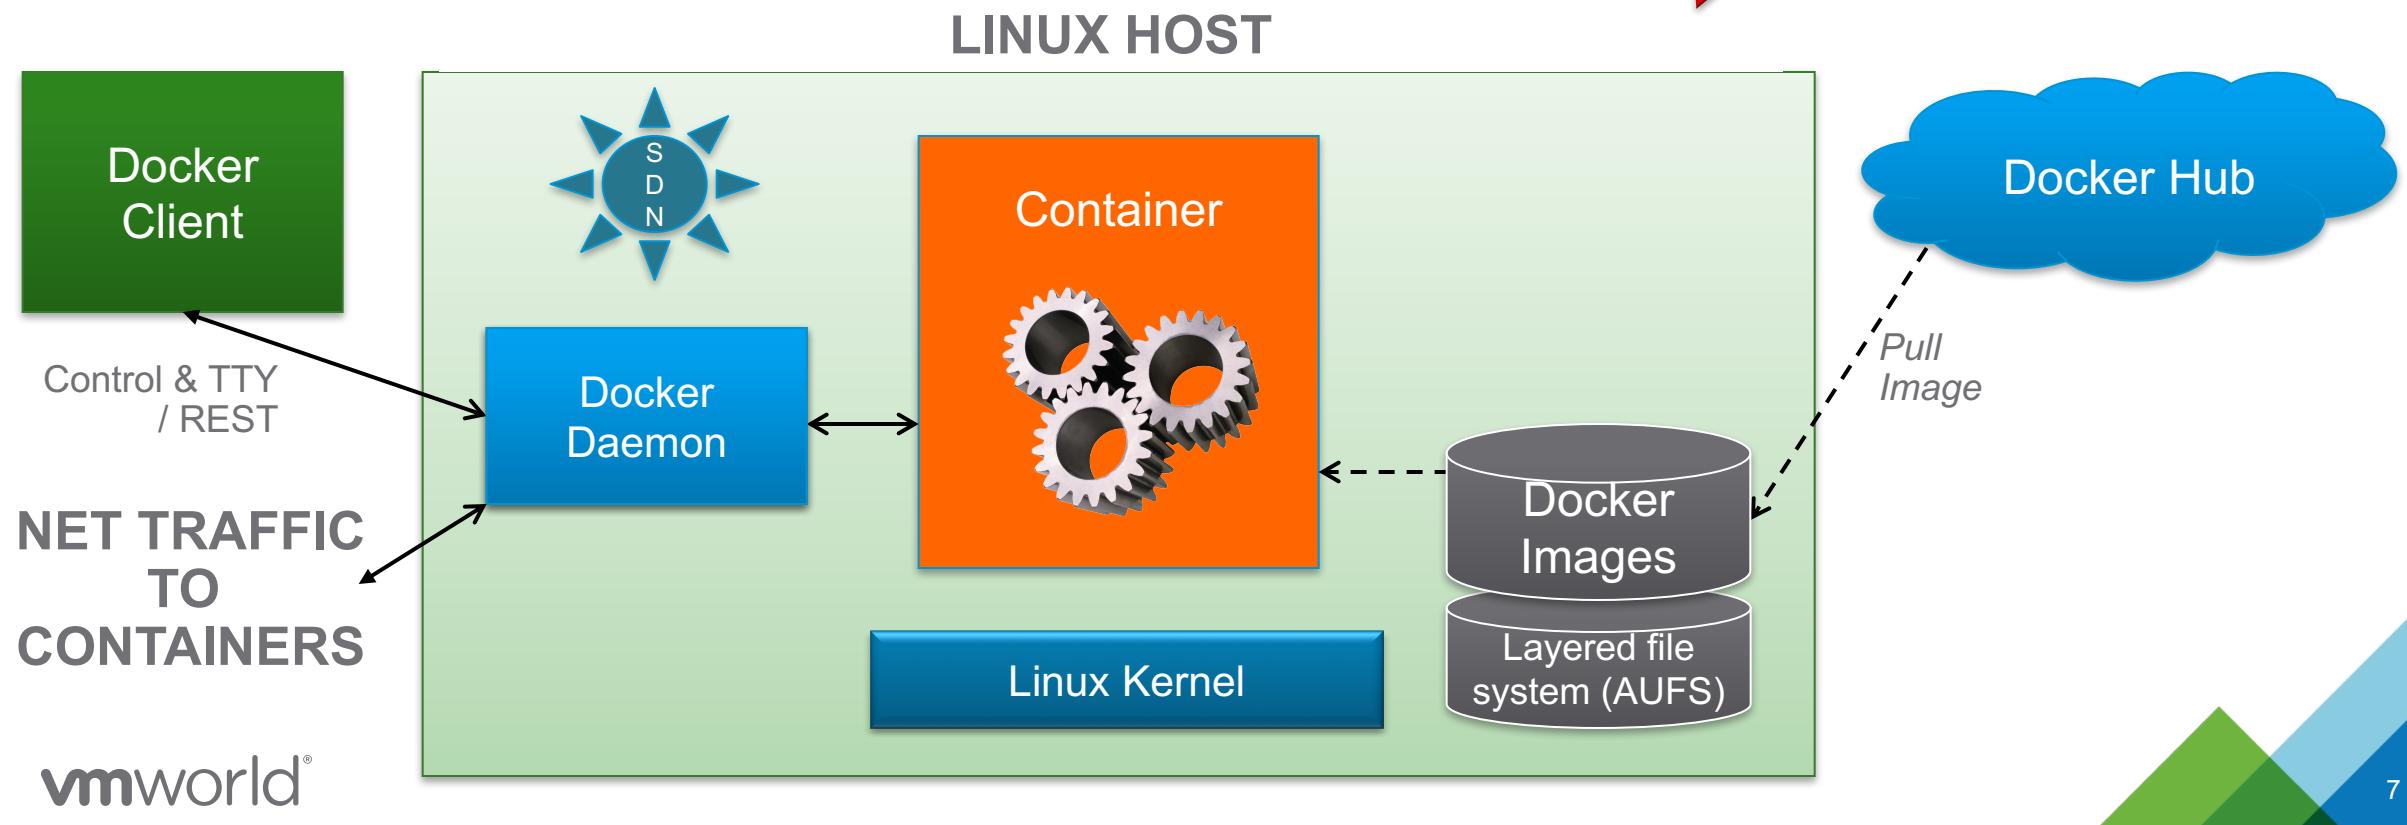
# Docker 101

```
vmware [ ~ ]$ sudo docker run -it ubuntu
root@636a8cb6d180:/# ls
bin  dev  home  lib64  mnt  proc  run  srv  tmp  var
boot etc  lib   media  opt  root  sbin  sys  usr
root@636a8cb6d180:/# █
```



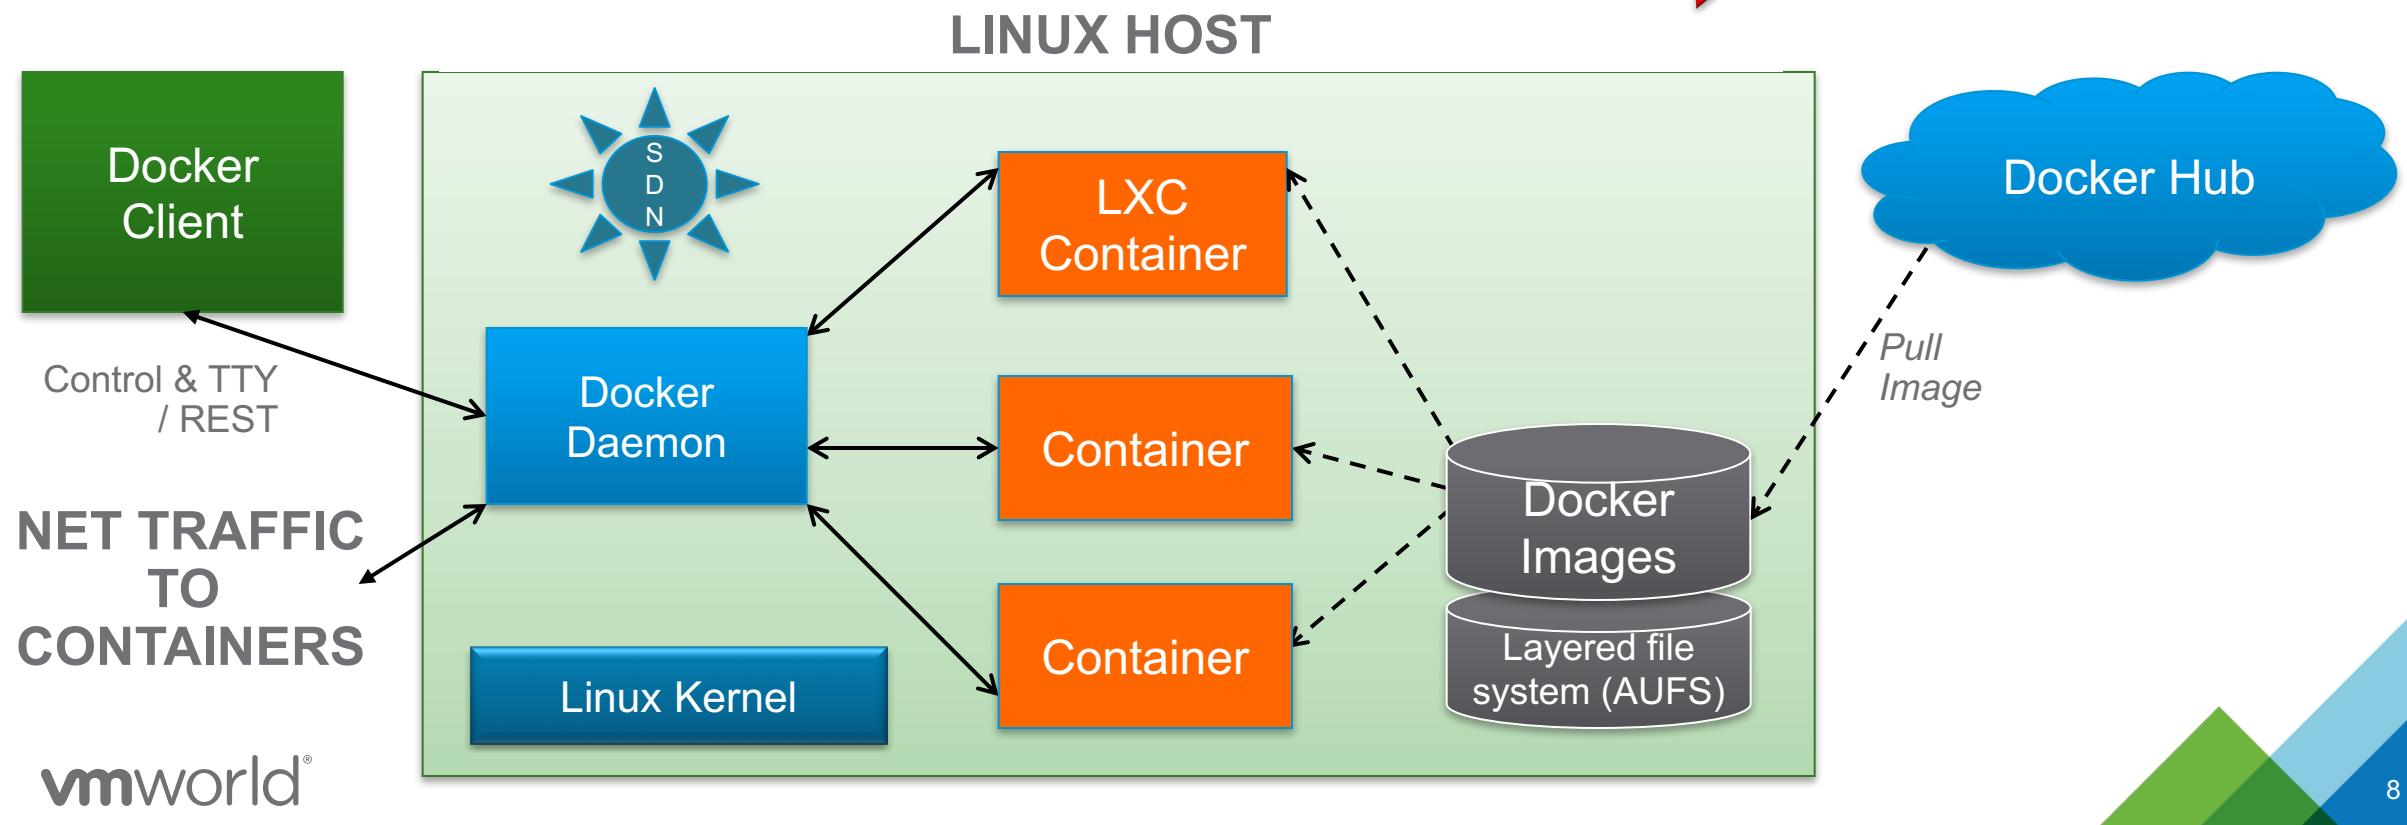
# What is a Container?

1. An executable process
2. Resource constraints / private namespace
3. Binary dependencies: Application, runtime, OS
4. A shared Linux kernel for running the executable
5. Ephemeral and persistent storage layers



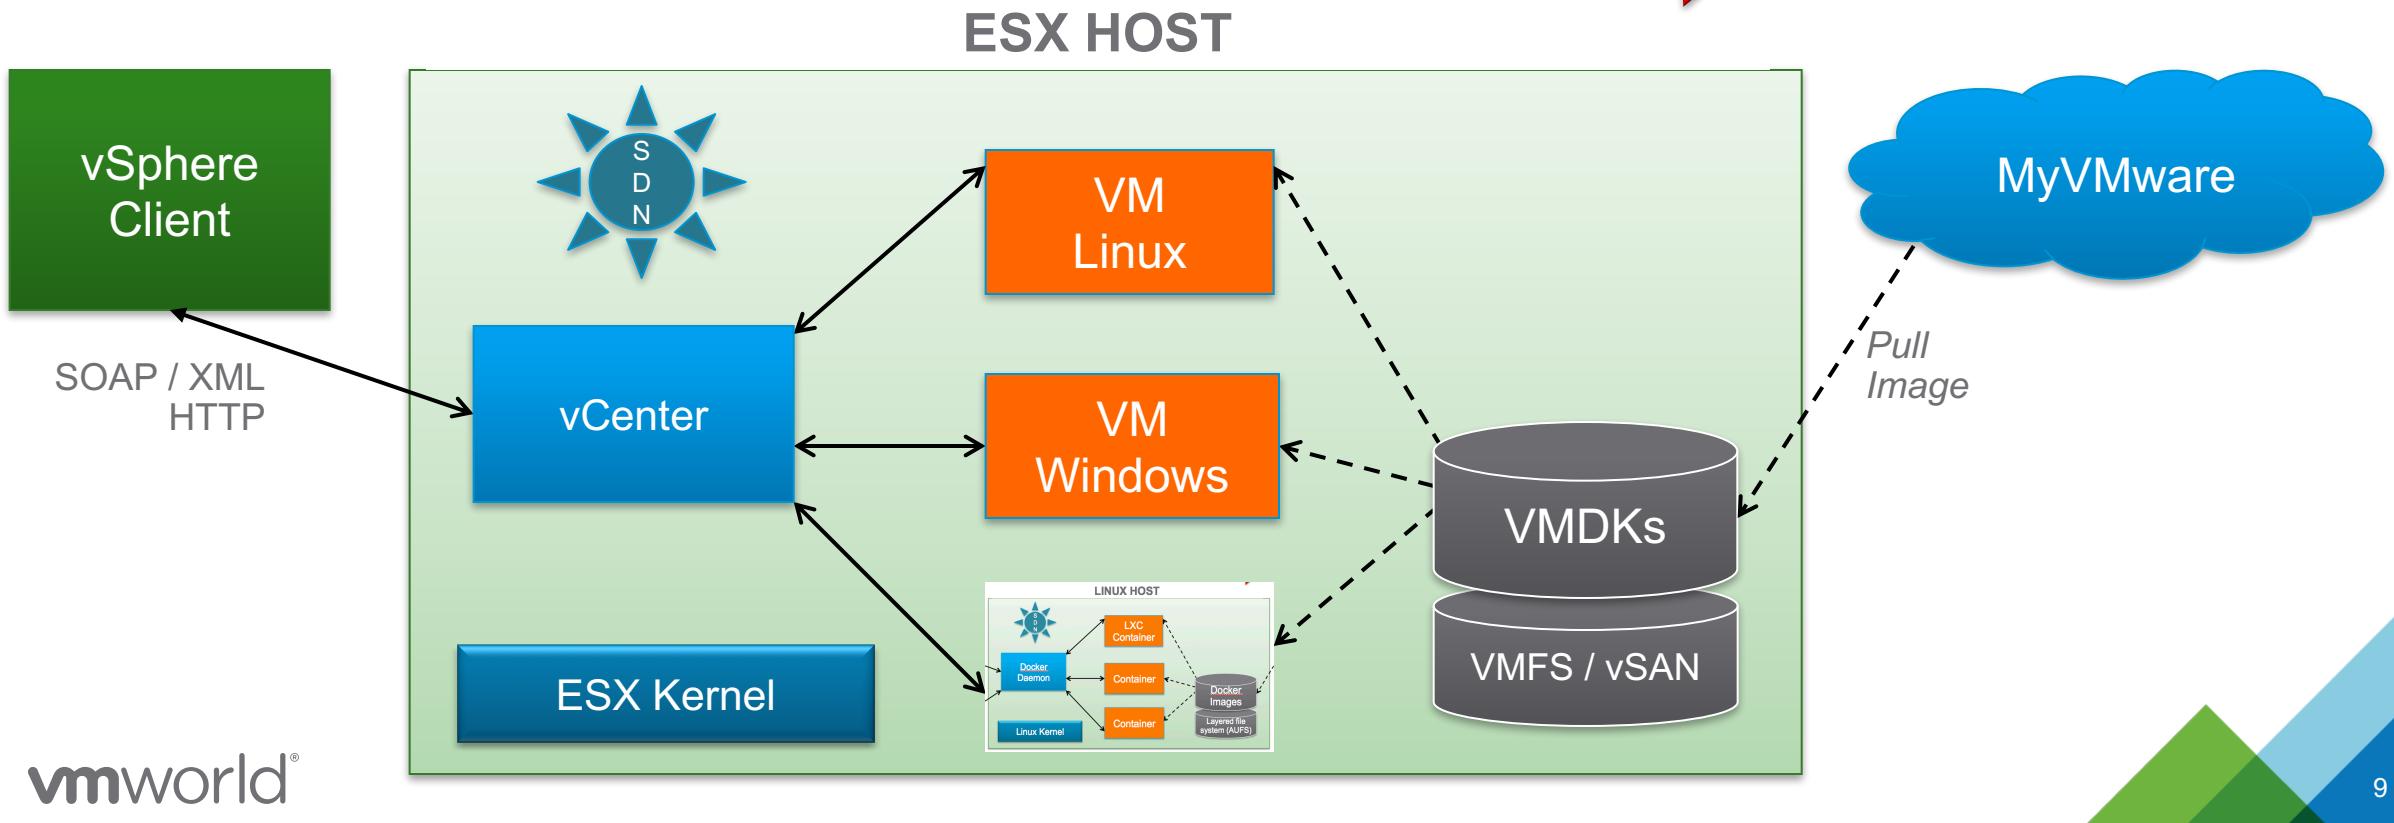
# What is a Container Host?

1. Control plane & lifecycle management for containers
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A single Linux kernel manages everything
5. A static size and a resource reservation when virtual



# What is a Hypervisor?

1. Control plane & lifecycle management for VMs
2. Resource scheduling and a VM abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A hypervisor kernel manages everything except apps
5. A static size, but no resource concerns unless nested



# Types of Containers

- ***Long-running***
  - Can be stateless or stateful
  - Eg. Application servers, databases, load-balancers, KV stores etc
  - Typically a need for strong isolation
- ***Transactional***
  - Runs for a period and transforms some data
  - Eg. Runs a build. Processes a web request. Batch processing
  - Should only consume resource when running
- ***Sidecar / helpers***
  - Augments the capabilities of a service or provides a helper function
  - Eg. Logging, monitoring, caching
  - Scales with the service. Potentially hindered by strong isolation

# Isolation Domains and Data Persistence

- How do I isolate workloads from each other?
  - Runtime isolation – resource limits, kernel panic, ESX host failure, rack or region failure
  - Network isolation – traffic sniffing, firewalls, encryption, rate limiting
  - Storage isolation – data persistence, backup, networking, RBAC
- Stateful vs Stateless / Cattle vs Pets
  - Different classifications of data. Where should it go?
    - Image state, container state, volume state. What's the difference?
  - Should data lifespan be inherently tied to compute (VM / container)?
  - Without live migration, shared storage and HA, stateless looks attractive
- The question is not “what plumbing do I need?”
  - What *characteristics* or business value do I need for my application?
  - Better question than, “do I need a container or a VM?”

# Characteristics: Containers vs VMs? So 2014!

- ***Speed***
  - Start time vs throughput. Benefits transactional containers
  - Hello World vs Tomcat. Less benefit for long-running apps
- ***Efficiency***
  - Less memory consumption? Depends when virtualized
  - Network traffic NAT'd through guest vs straight to vNIC
- ***Portability***
  - Docker image abstraction and API is very portable
  - Subtle issues with kernel versions and patches
- ***Isolation***
  - Shared kernel significantly reduces runtime isolation
- ***Granularity***
  - Containers are great for granular services. Group in a VM as isolation domain.

# Container Workflow Efficiencies

- **A Portable Runtime Platform**
  - “Package once, run anywhere”
  - Defacto runtime model for higher level frameworks
- **The Container Model**
  - “Docker is to apt what apt is to tar” – dependency management
  - Snapshotting, image format, Docker Hub, state management (volumes etc)
  - Scripting and automation – express an environment in a text file
  - Predictable initialization state – predictability vs. reconfiguration
- **Continuous Integration**
  - Secure registries, integration with popular tooling
  - Rapid improvements in tooling, monitoring, log integration etc.

# The Container Stack and Its Challenges

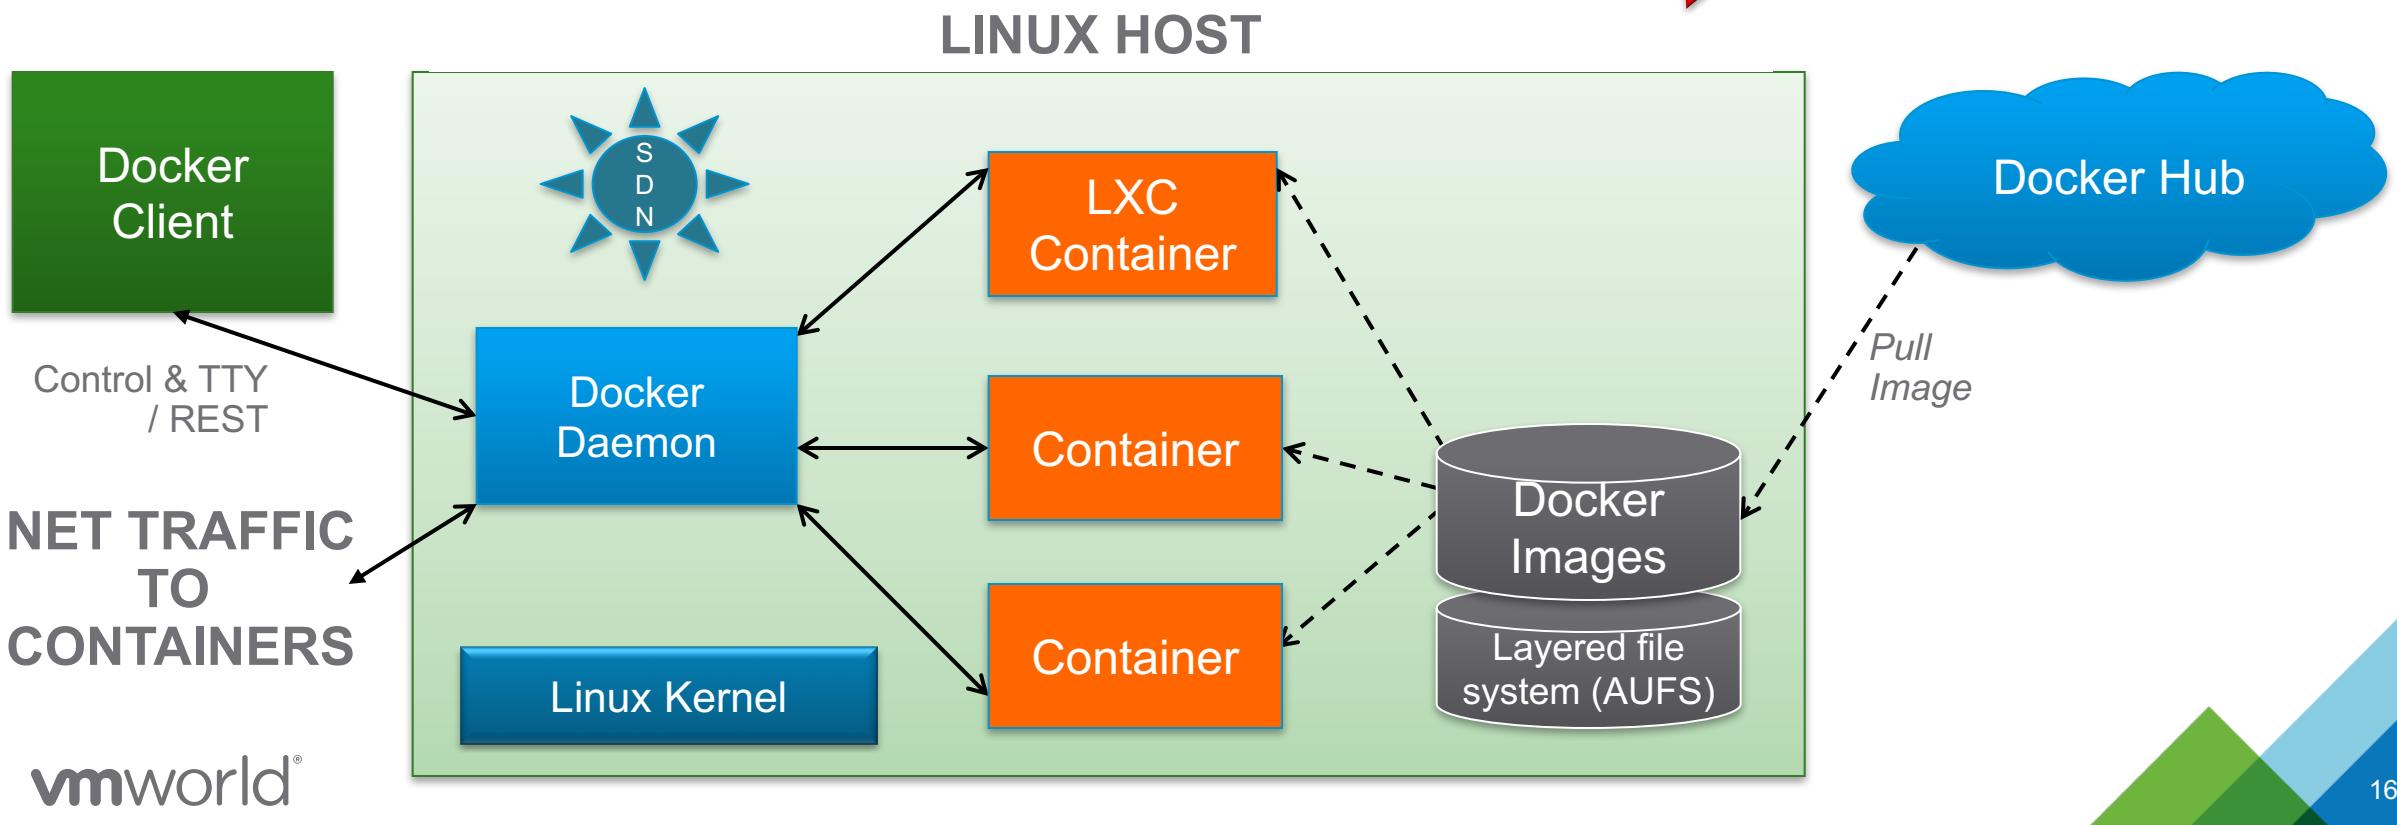
- I want to run containers. Where do I start?
- Should I run containers alongside my other apps or create a silo?
- Should I be considering bare metal?
  - So how many containers should I have in each VM?
  - How big should each container host be? Impact of re-configuration?
  - Which containers can be safely co-located with which other containers?
  - Which Linux distro do I want and how to handle the high patch cadence?
  - Who is responsible for infrastructure admin and how does it integrate?

# So What is VIC and How does it help?

- **VIC brings all of the container workflow efficiencies to vSphere infrastructure**
  - Control vSphere infrastructure from a Docker client without having any vSphere credentials
  - Treat VMs as ephemeral containers
  - No more OVAs, VMDKs, Templates, Cloning. Push / pull your state from secure registries
- **VIC allows you to translate business value into plumbing**
  - Eg. I need to deploy Wordpress with MySQL.
  - Do I want strong isolation between these workloads? Container as a VM
  - Do I want strong isolation from other tenants? Container *in* a VM
- **VIC helps to draw clear lines between admins and users**

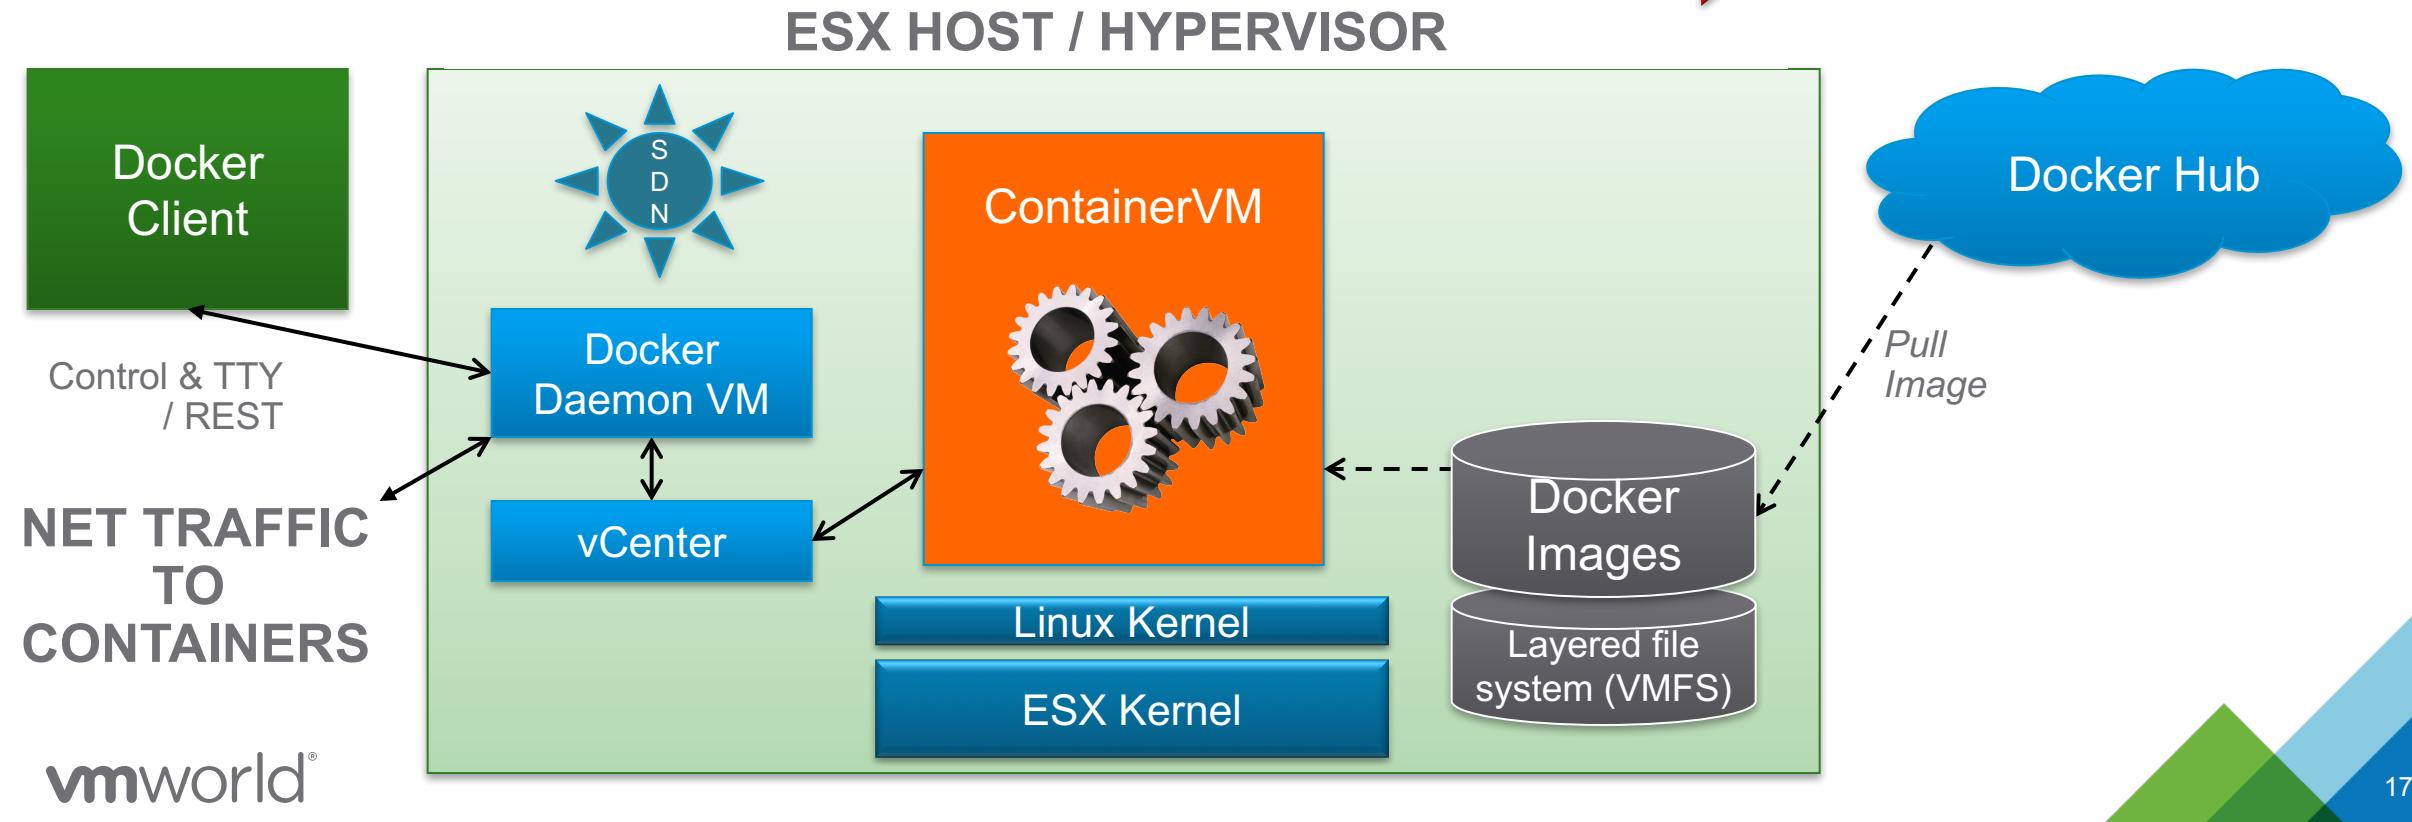
# Revisit: What is a Container Host?

1. Control plane & lifecycle management for containers
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A single Linux kernel manages everything
5. A static size and a resource reservation when virtual



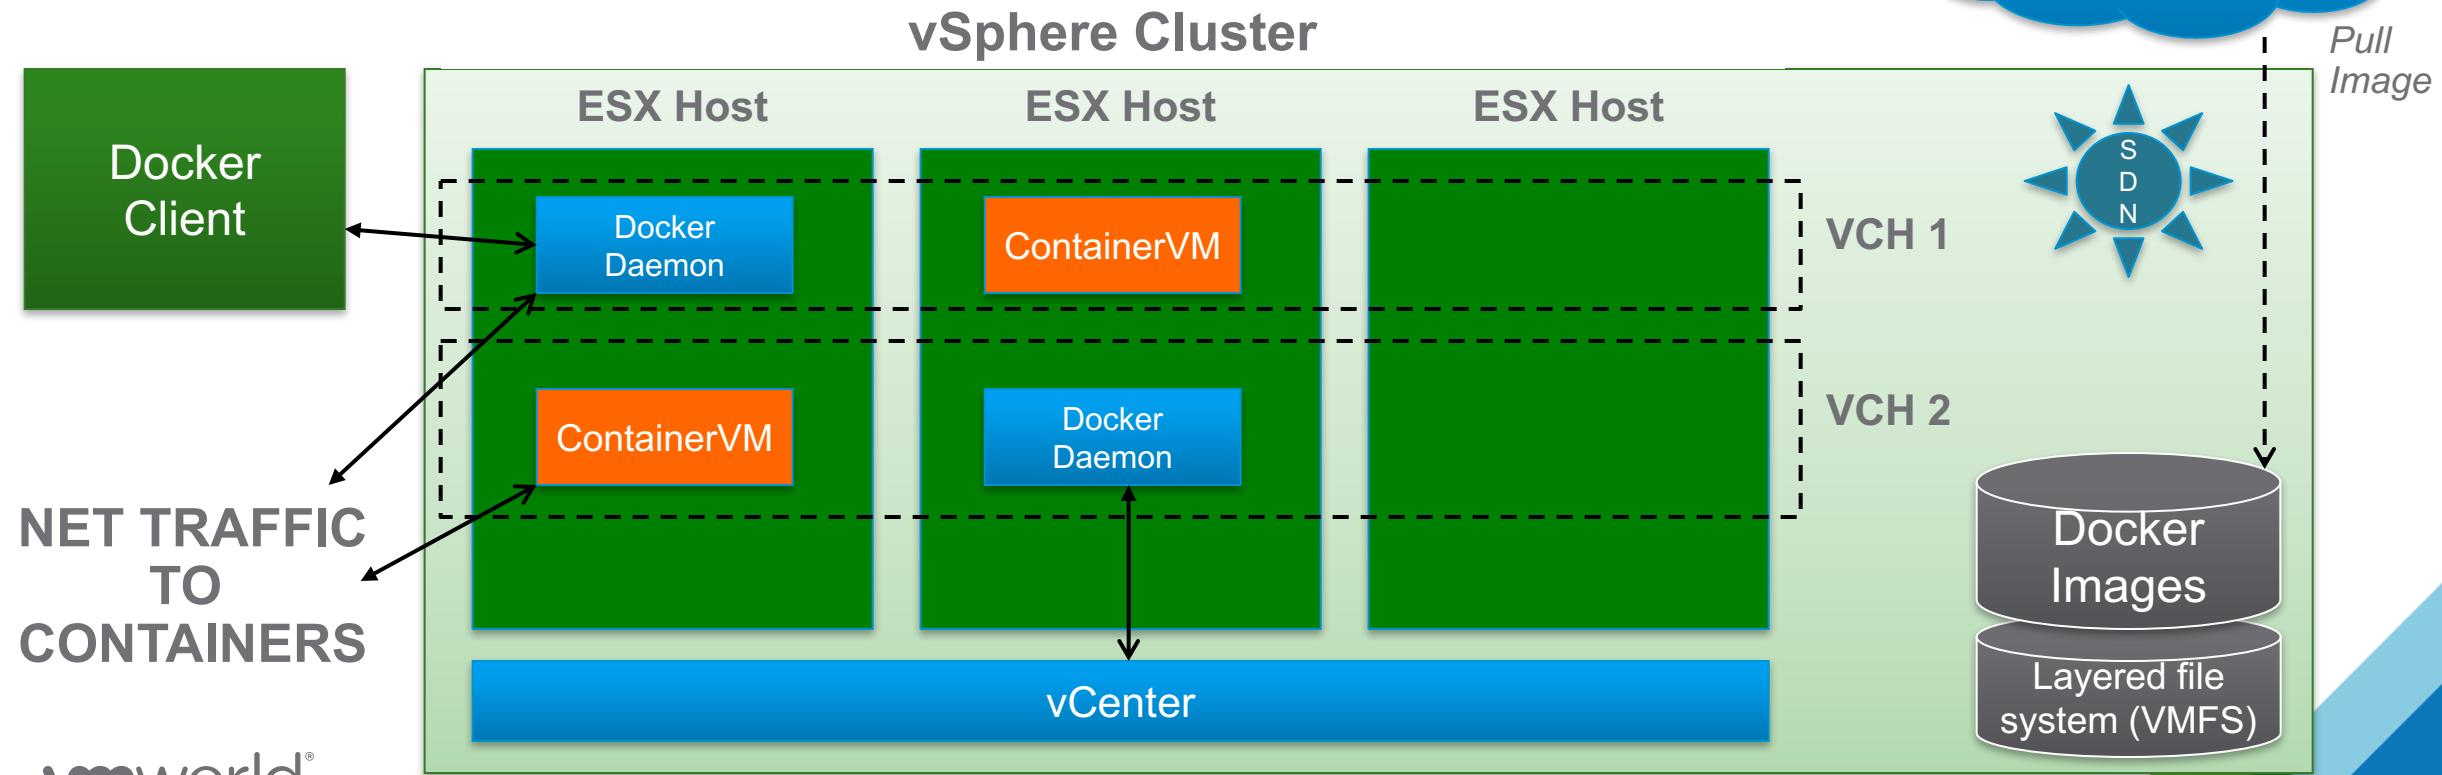
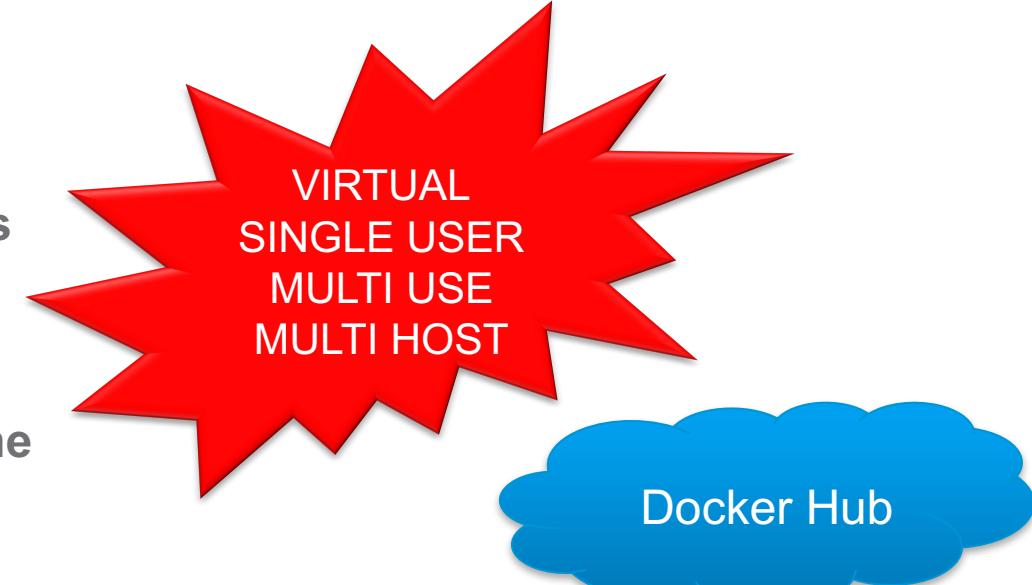
# What is a “ContainerVM”?

1. An executable process
2. Resource constraints / private namespace
3. Binary dependencies: Application, runtime, OS
4. A *private* Linux kernel for running the executable
5. Ephemeral and persistent storage layers



# What is a Virtual Container Host?

1. Control plane & lifecycle management for ContainerVMs
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A Linux kernel per container, separate from control plane
5. Dynamic size and a resource *limit*, not reservation!



# VIC Roadmap

- **VIC 1.0**
  - Docker engine and registry (Harbor) shipped in Q4 2016
  - MVP Docker commands. Fundamentals of integration
  - vSAN, vMotion, DRS, NSX distributed port groups
- **VIC 1.1**
  - Additional Docker capabilities – Eg. compose, exec
  - Critical bug fixes
  - Deeper SDDC integrations – Eg. HA
- **VIC 1.2**
  - Security enhancements – Image signing, vulnerability scanning
  - Additional Docker capabilities
  - Deeper NSX integration
  - Further SDDC integrations – vVols, SDRS

## Questions

@bensdoings

[Github.com/vmware/vic-product](https://github.com/vmware/vic-product)

# Demo - VIC

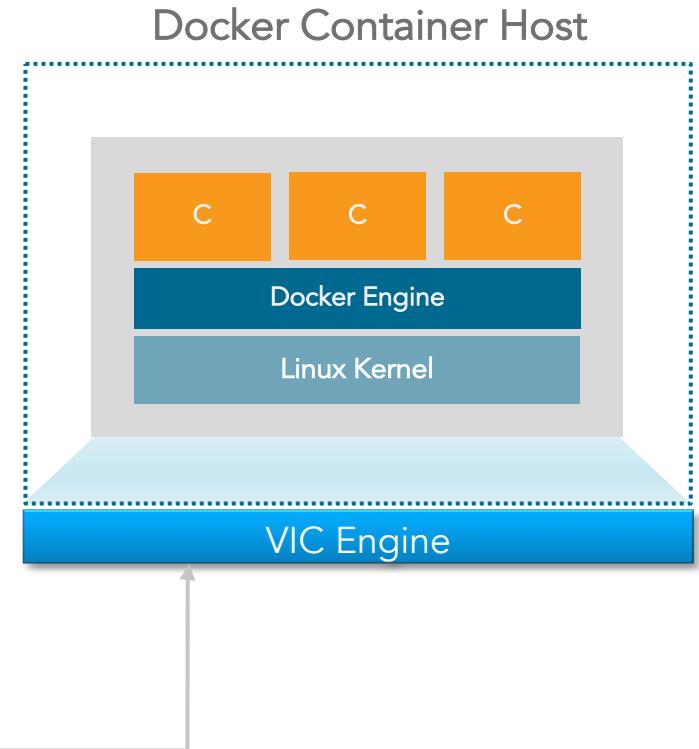
# New in 1.2 - Developer Sandbox

- Developer self-service with VI Admin governance
  - Developer consumes resources via Docker API/CLI
- Provides developers with self-service for applications not yet in the enterprise service catalog
  - Rapid prototyping
- Run a full-fledged docker engine as a ContainerVM using vSphere Integrated Containers and the Docker API/CLI



Scott : Developer  
Write & deploy code

```
docker run -p 12375:2375 -d vmware/dch-photon
```



# Project Harbor

# Harbor

---

**1** Container Image Basics

---

**2** Project Harbor Introduction

---

**3** Consistency of Images

---

**4** Security

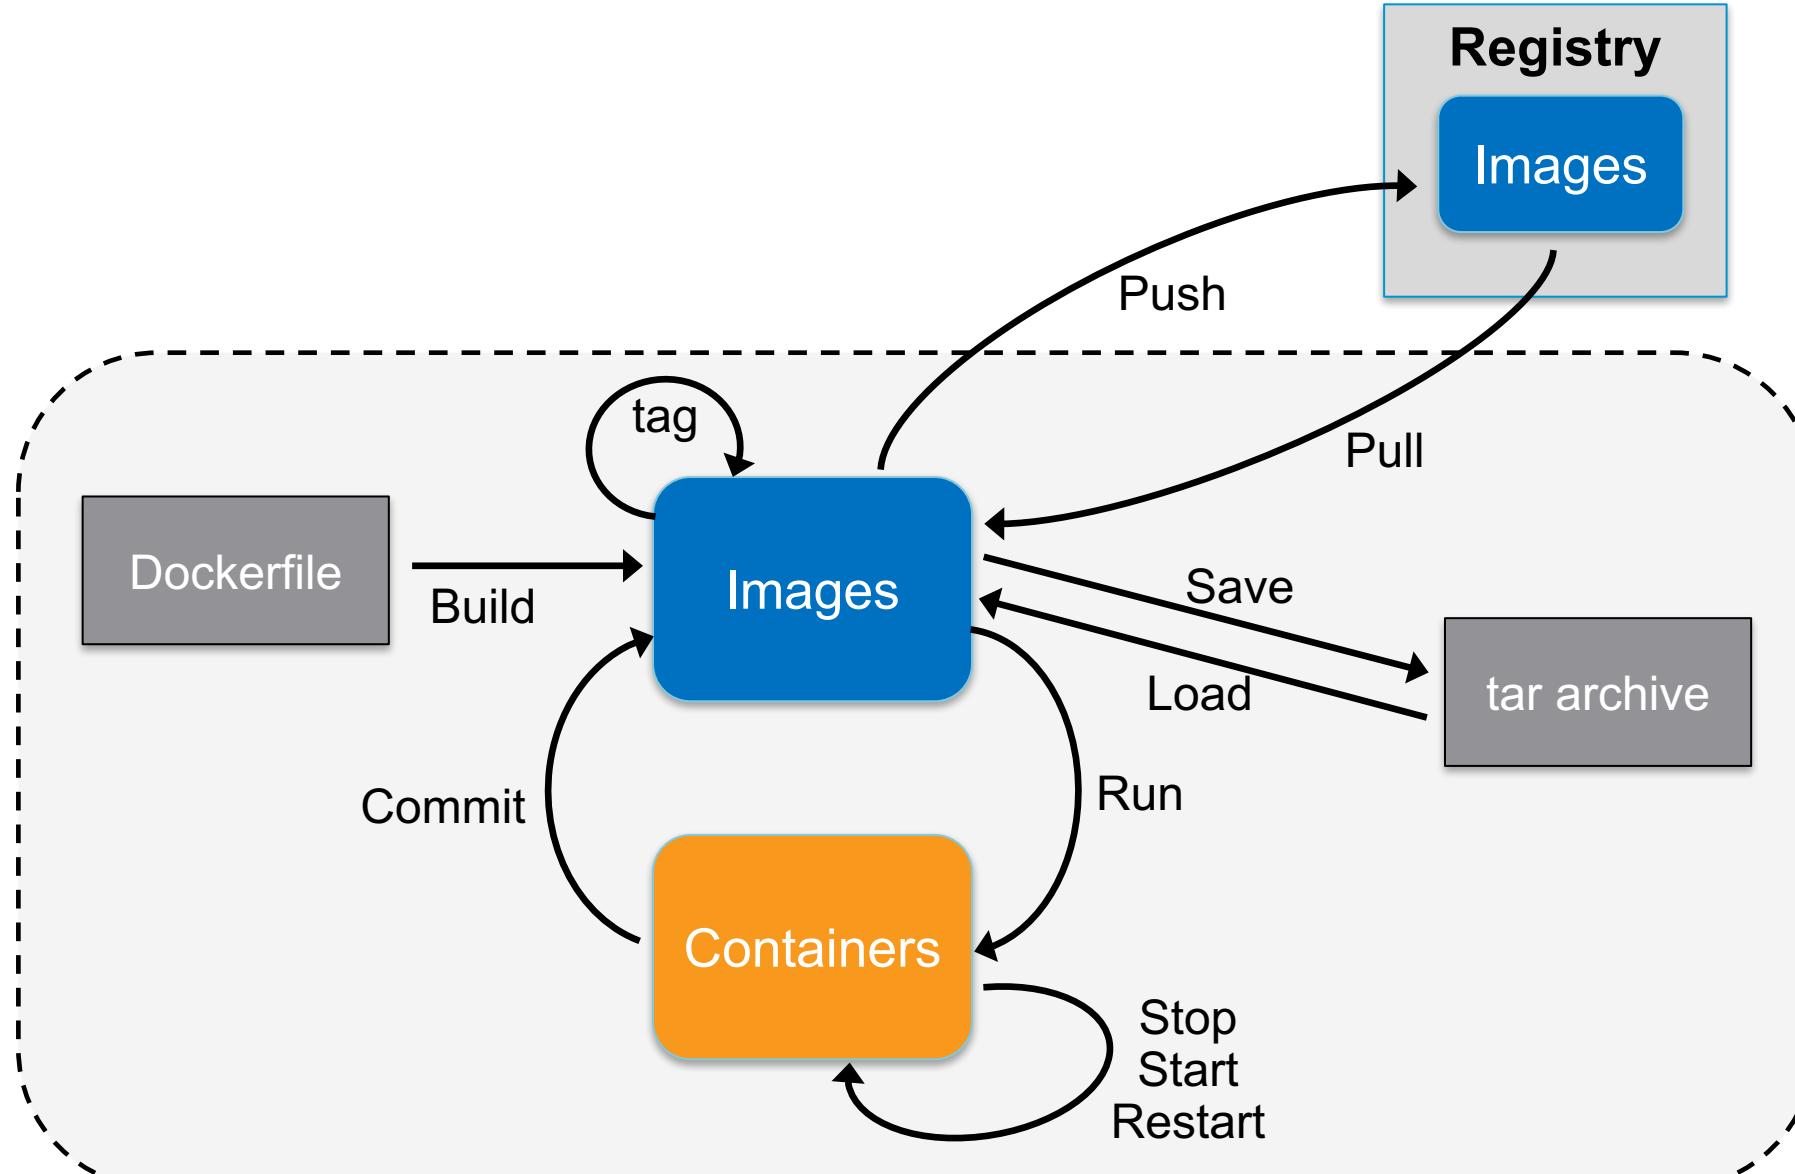
---

**5** Image Distribution

---

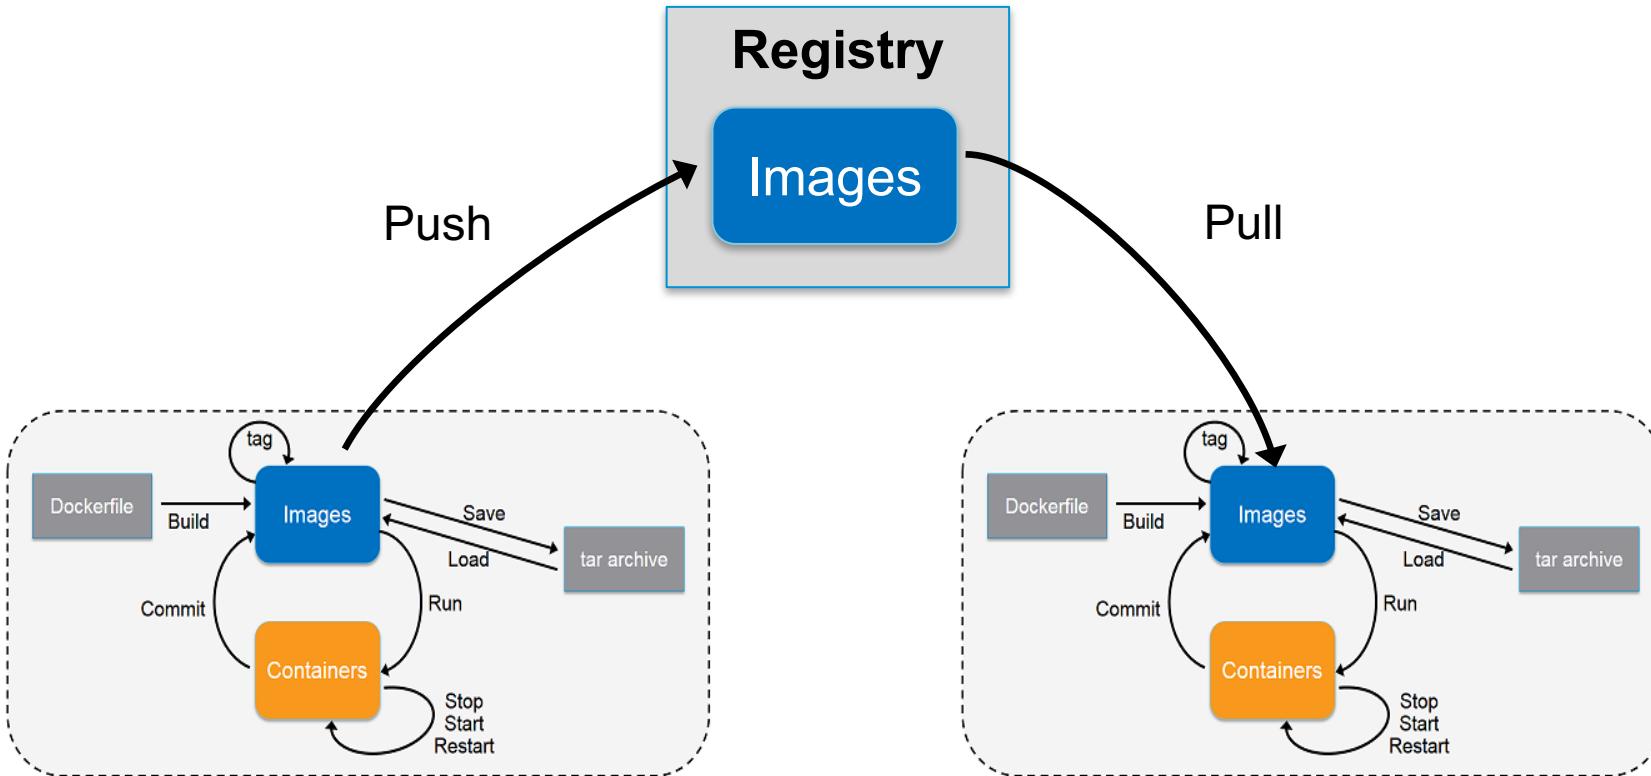


# Lifecycle of Containers and Images



# Registry - Key Component to Manage Images

- Repository for storing images
- Intermediary for shipping and distributing images
- Ideal for access control and other image management



# Harbor

---

- [\*\*1\*\* Container Image Basics](#)
- [\*\*2\*\* Project Harbor Introduction](#)
- [\*\*3\*\* Consistency of Images](#)
- [\*\*4\*\* Security](#)
- [\*\*5\*\* Image Distribution](#)
- [\*\*6\*\* High Availability of Registry](#)

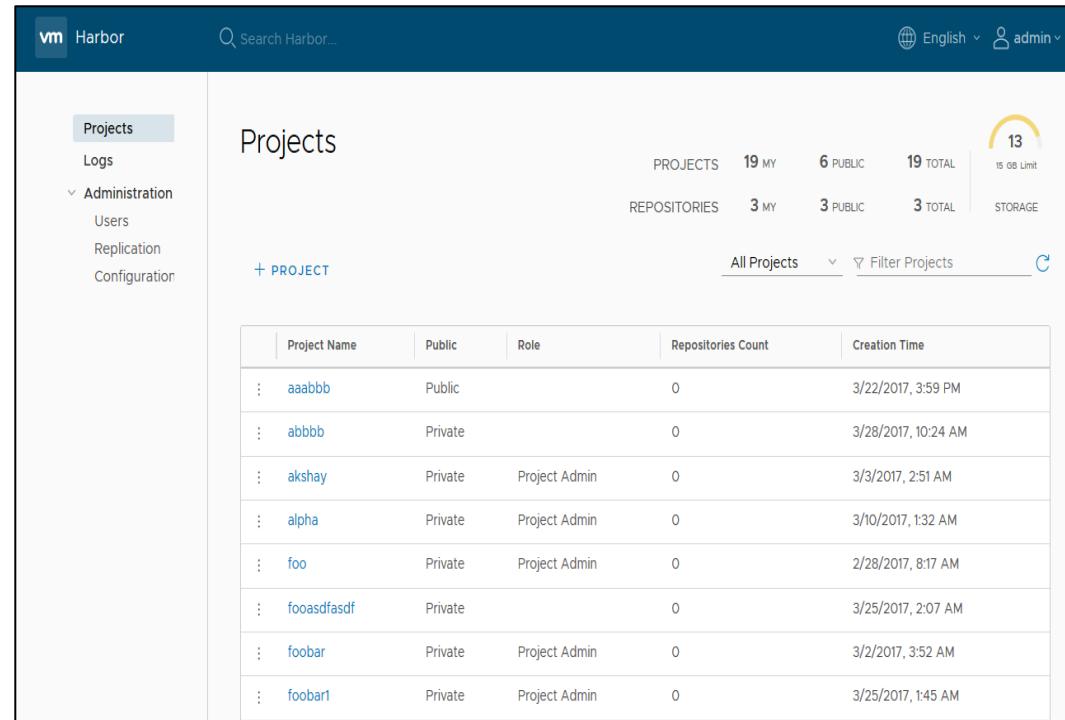
# Project Harbor



- An open source enterprise-class registry server.
- Initiated by VMware China, adopted by users worldwide.
- Integrated into vSphere Integrated Containers.
- Apache 2 license.
- <https://github.com/vmware/harbor/>

# Key Features

- User management & access control
  - RBAC: admin, developer, guest
  - AD/LDAP integration
- Policy based image replication
- Notary
- Vulnerability Scanning
- Web UI
- Audit and logs
- Restful API for integration
- Lightweight and easy deployment



The screenshot shows the Harbor web interface with a dark blue header. On the left is a sidebar with 'Projects' selected, followed by 'Logs', 'Administration' (with 'Users', 'Replication', and 'Configuration' sub-options), and a '+ PROJECT' button. The main area is titled 'Projects' and displays summary statistics: 19 PROJECTS (19 MY, 6 PUBLIC, 19 TOTAL), 3 REPOSITORIES (3 MY, 3 PUBLIC, 3 TOTAL), and 13 STORAGE (15 GB Limit). Below this is a table of projects:

Project Name	Public	Role	Repositories Count	Creation Time
aaabbb	Public		0	3/22/2017, 3:59 PM
abbbb	Private		0	3/28/2017, 10:24 AM
akshay	Private	Project Admin	0	3/3/2017, 2:51 AM
alpha	Private	Project Admin	0	3/10/2017, 1:32 AM
foo	Private	Project Admin	0	2/28/2017, 8:17 AM
foasdfasdf	Private		0	3/25/2017, 2:07 AM
foobar	Private	Project Admin	0	3/2/2017, 3:52 AM
foobar1	Private	Project Admin	0	3/25/2017, 1:45 AM

# Users and Developers

- Users



Downloads



Stars



Users

- Developers



Forks



Contributors

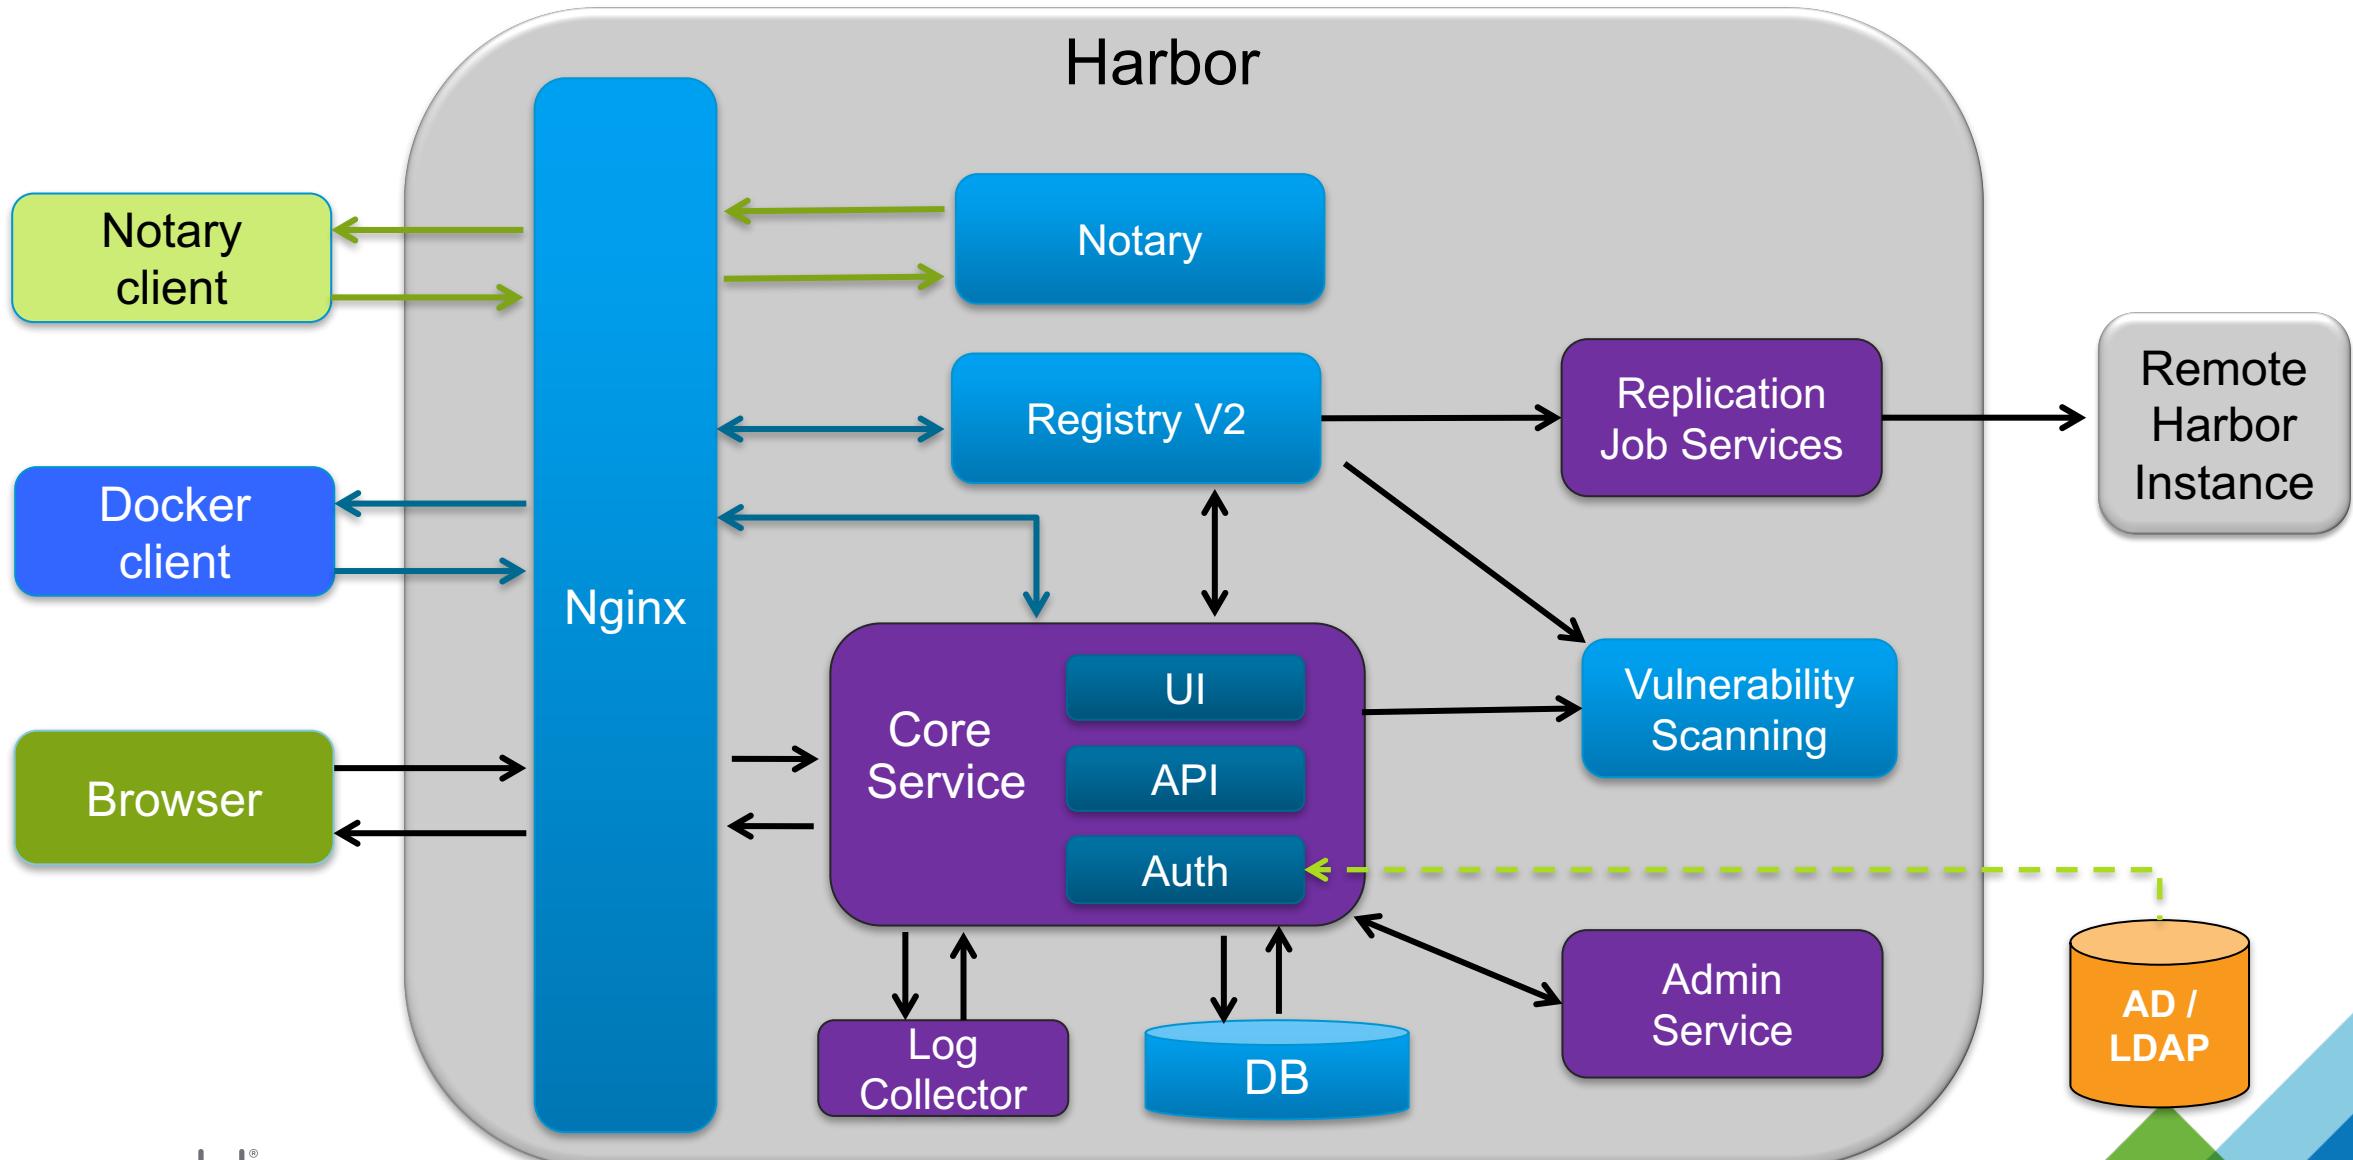


Partners

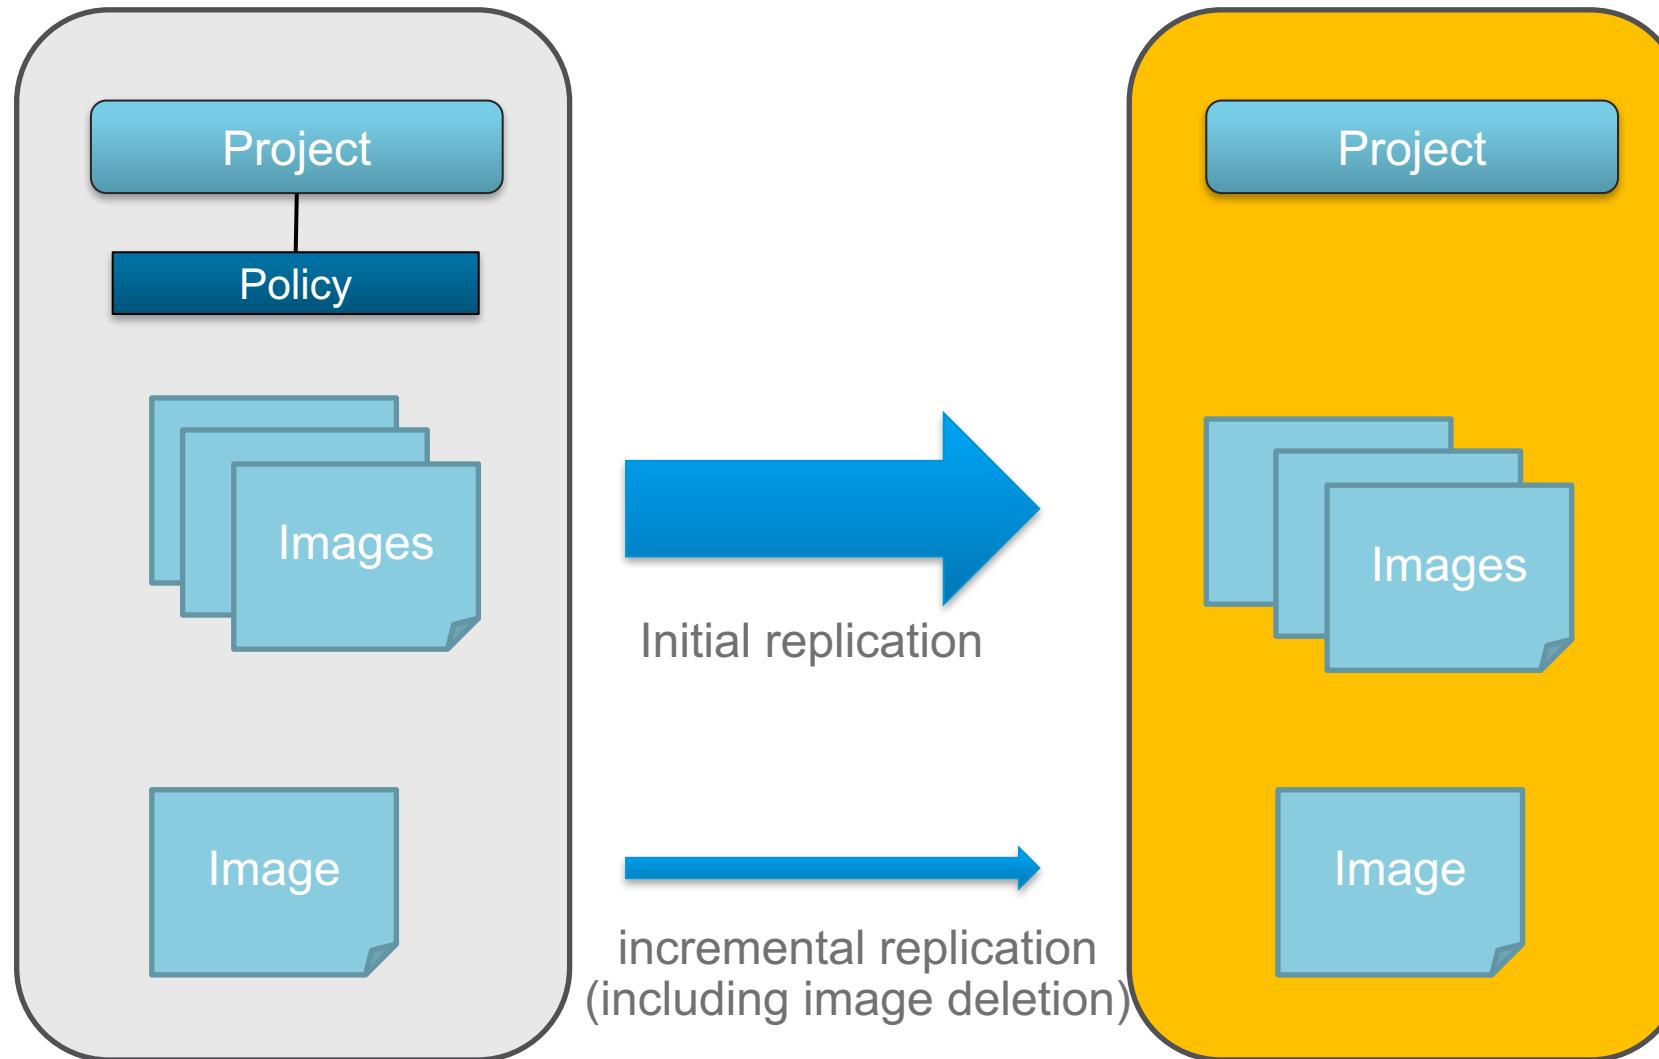
# Harbor users and partners



# Harbor Architecture



# Image replication (synchronization)



# Harbor

---

1 Container Image Basics

---

2 Project Harbor Introduction

---

3 Consistency of Images

---

4 Security

---

5 Image Distribution

---



# Consistency of Container Images

- Container images are used throughout the life cycle of software development
  - Dev
  - Test
  - Staging
  - Production
- Consistency must be maintained
  - Version control
  - Issue tracking
  - Troubleshooting
  - Auditing

# Same Dockerfile Always Builds Same Image?

Example:

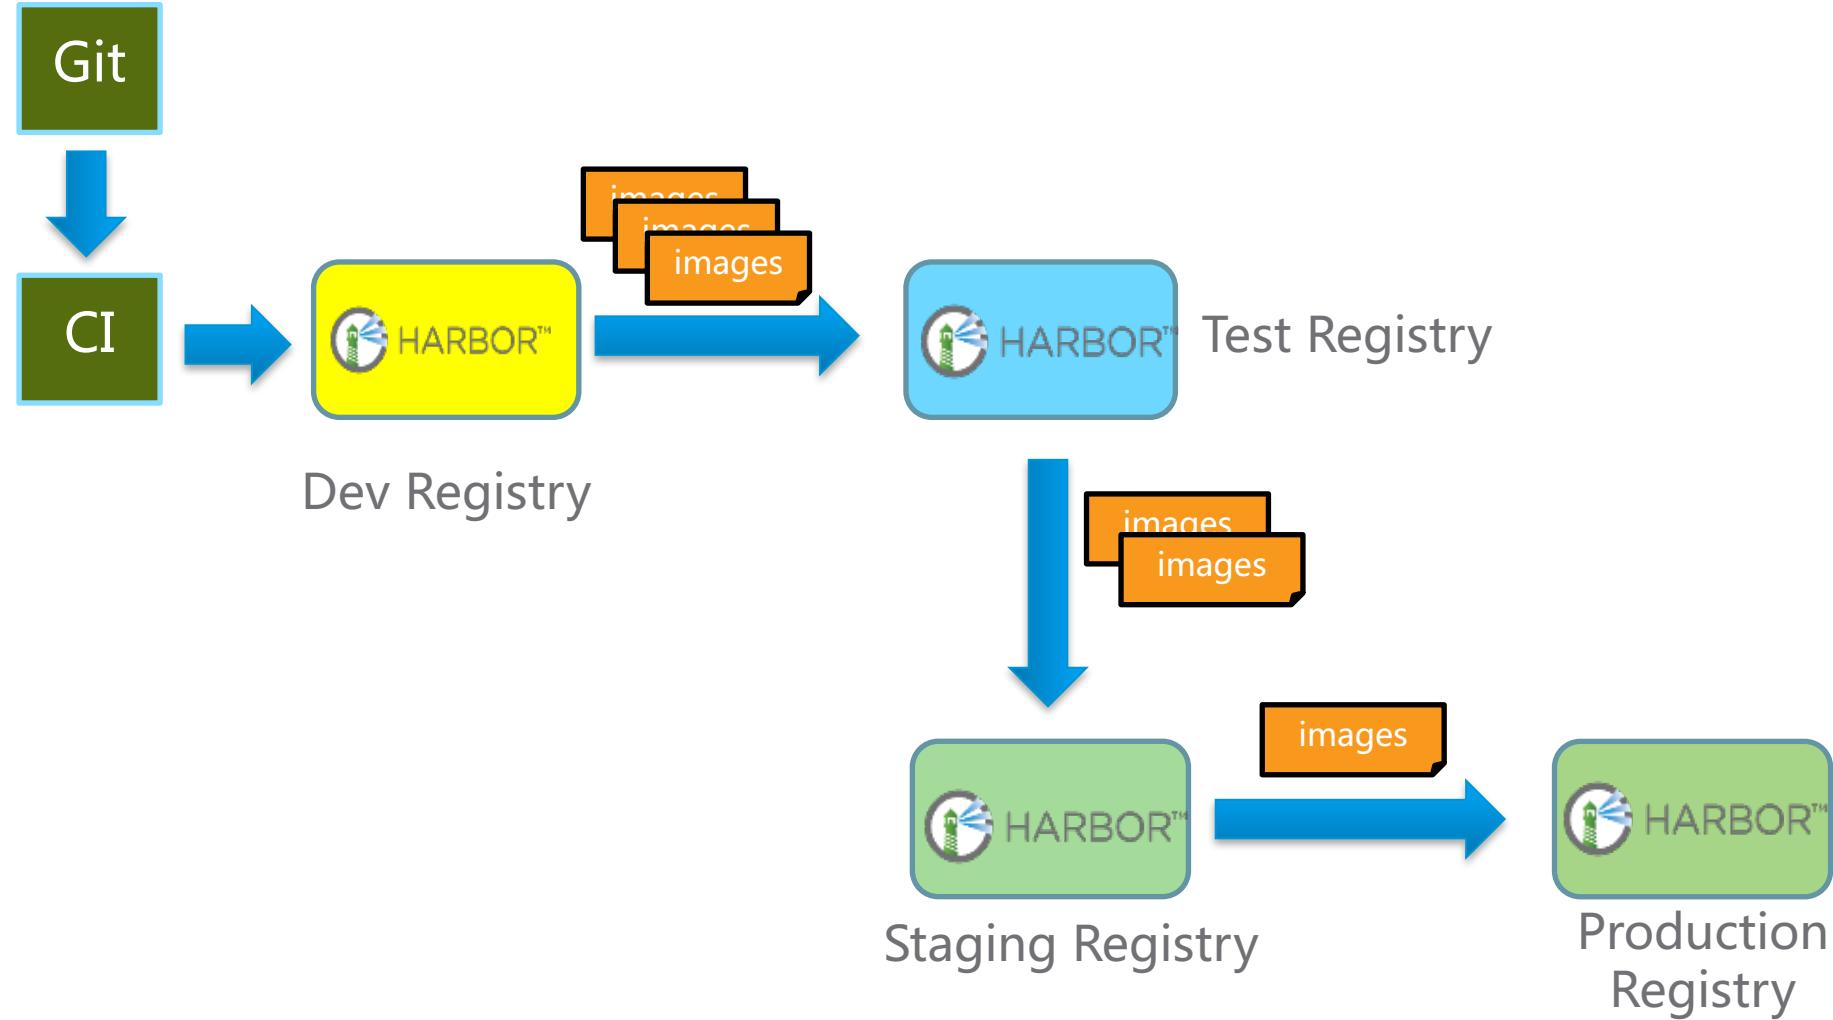
```
FROM ubuntu

RUN apt-get install -y python

ADD app.jar /myapp/app.jar
```

- Base image `ubuntu:latest` could be changed between builds
- `ubuntu:14.04` could also be changed due to patching
- `apt-get` (`curl`, `wget..`) cannot guarantee always to install the same packages
- `ADD` depends on the build time environment to add files

# Shipping Images in Binary Format for Consistency



Images are synchronized between environments  
by using Harbor registry.

# Harbor

---

**1** Container Image Basics

---

**2** Project Harbor Introduction

---

**3** Consistency of Images

---

**4** Security

---

**5** Image Distribution

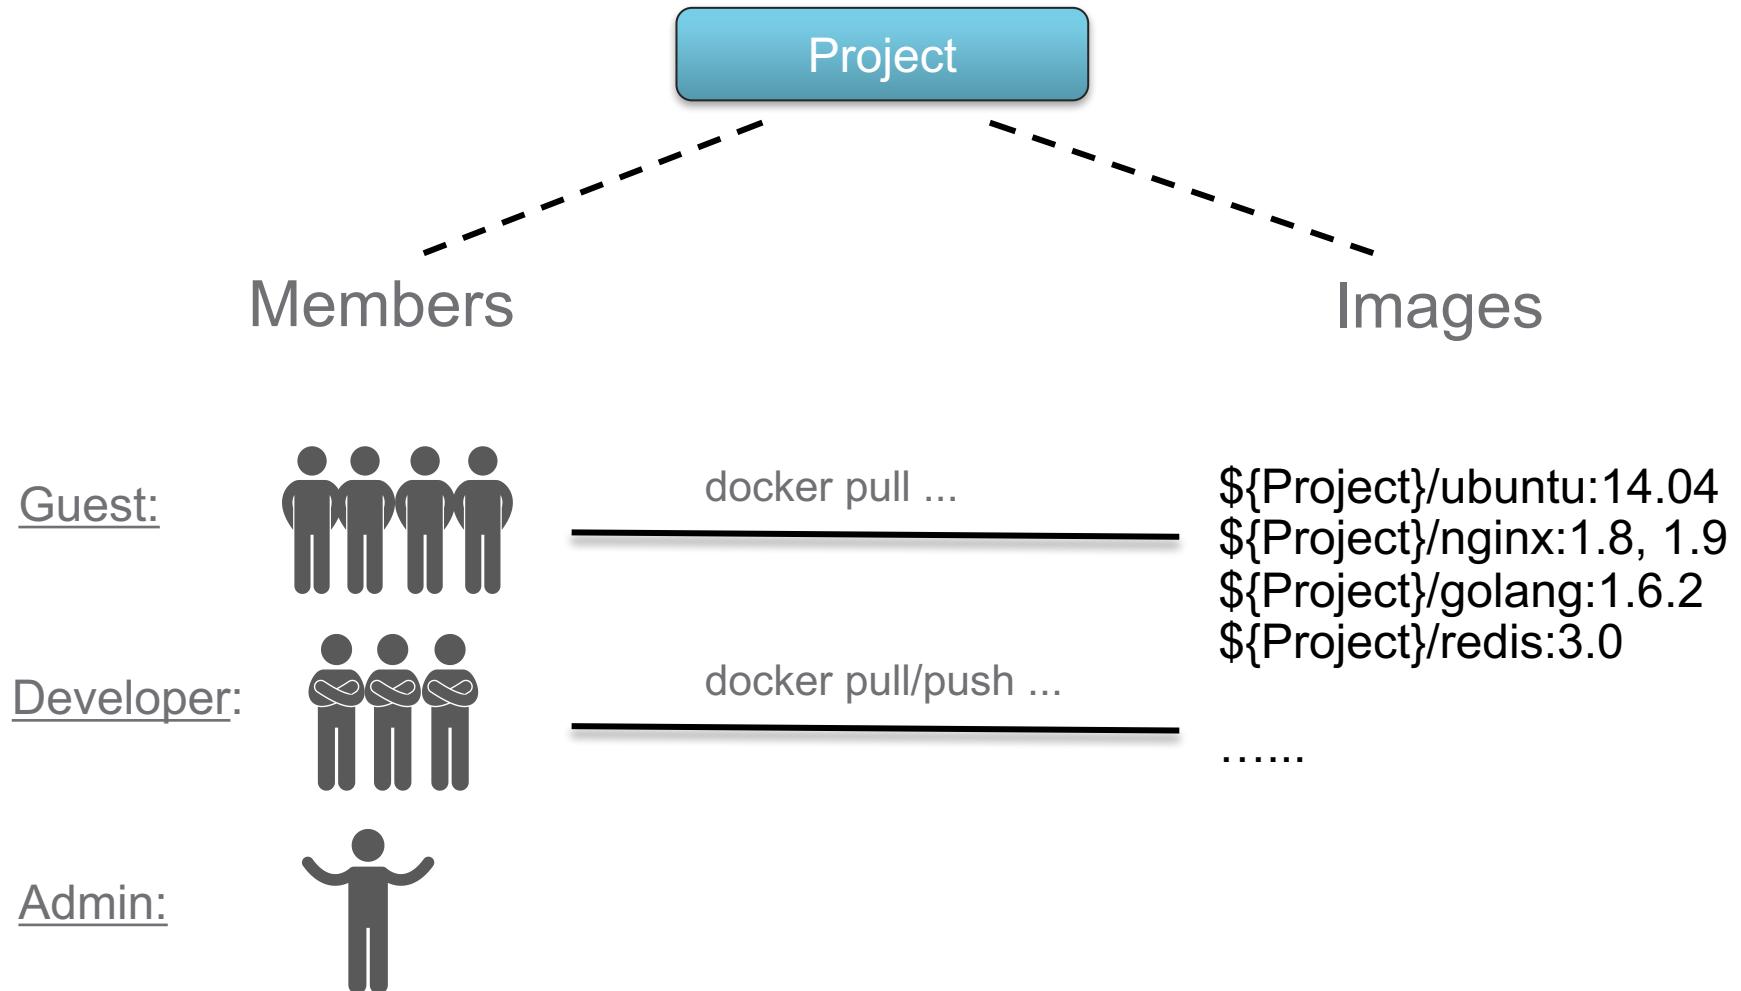
---



# Access Control to Images

- Organizations often keep images within their own organizations
  - Intellectual property stays in organization
  - Efficiency: LAN vs WAN
- People with different roles should have different access
  - Developer – Read/Write
  - Tester – Read Only
- Different rules should be enforced in different environments
  - Dev/test env – many people can access
  - Production – a limited number of people can access
- Can be integrated with internal user management system
  - LDAP/Active Directory

# Example: Role Based Access Control in Harbor



# Other security considerations

- Enable content trust by installing Notary service
  - Image is signed by publisher's private key during pushing
  - Image is pulled using digest
- Perform vulnerability scanning
  - Identify images with vulnerabilities during pushing
  - Prevent images with vulnerabilities from being pulled
  - Regular scanning based on updated vulnerability database

# Harbor

---

1 Container Image Basics

---

2 Project Harbor Introduction

---

3 Consistency of Images

---

4 Security

---

5 Image Distribution

---

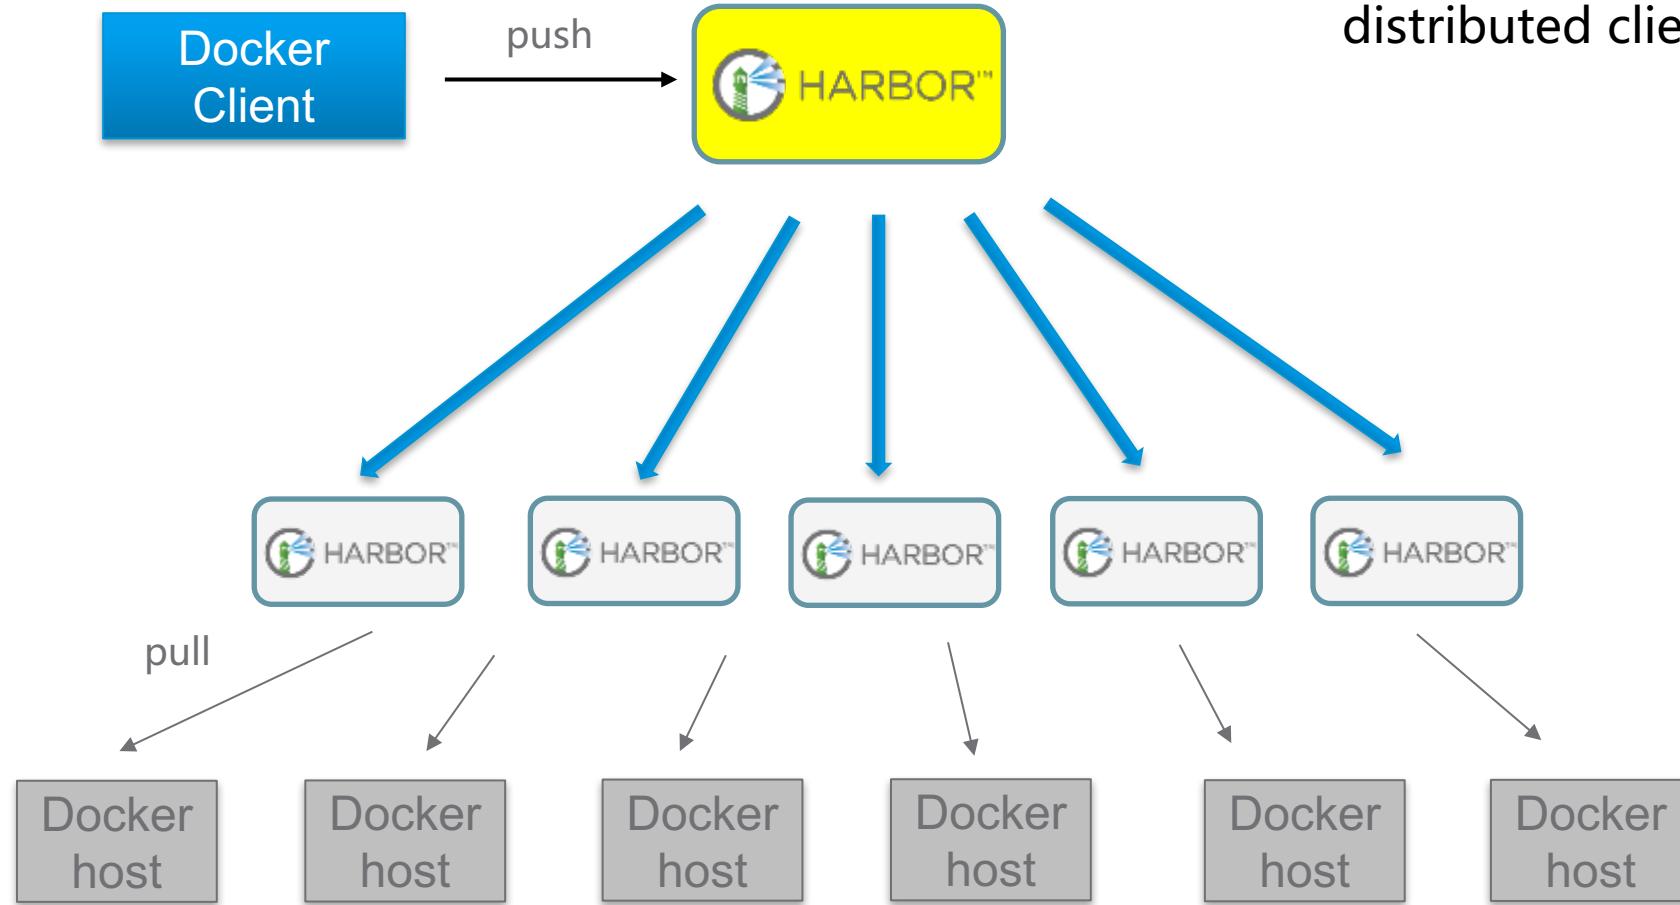


# Image Distribution

- Container images are usually distributed from a registry.
- Registry becomes the bottleneck for a large cluster of nodes
  - I/O
  - Network
- Scaling out an registry server
  - Multiple instances of registry sharing same storage
  - Multiple instances of independent registry sharing no storage

# Image Distribution via Master-Slave Replication

- Load balancing
- Works well with geographically distributed clients



# Demo: Harbor Registry

# Admiral Overview



## What is Admiral?

- A lightweight **Container Management Platform**
- Deploy and manage container based applications.
- Provide placement based on dynamic policy allocation.
- Container Management Layer in both vRA and VIC.
- Works with VIC engine or Docker host, Harbor registry.

# Container Deployment

Provision container with the image from the specified registry



ADMIRAL™

The screenshot shows the vSphere Integrated Containers interface. The left sidebar is titled "PROJECT" and shows "vmworld-group1". Under "Deployments", "Containers" is selected. The main area is titled "Containers (2)" and shows two running containers:

- myVmworld-04-mcm45...** (Running for 20 hours) - Host: 10.162.6.125, Created: August 17, 2017 6:25 PM, Ports: http://10.162.6.125:80:62017/tcp, Command: [redacted]
- vmworld-mcm50-513756...** (Running for 15 hours) - Host: 10.162.6.125, Created: August 17, 2017 11:00 PM, Ports: http://10.162.6.125:8088:62017/tcp, Command: [redacted]

The screenshot shows the "Provision a Container" dialog box. The left sidebar is identical to the previous screenshot. The main form has the following fields:

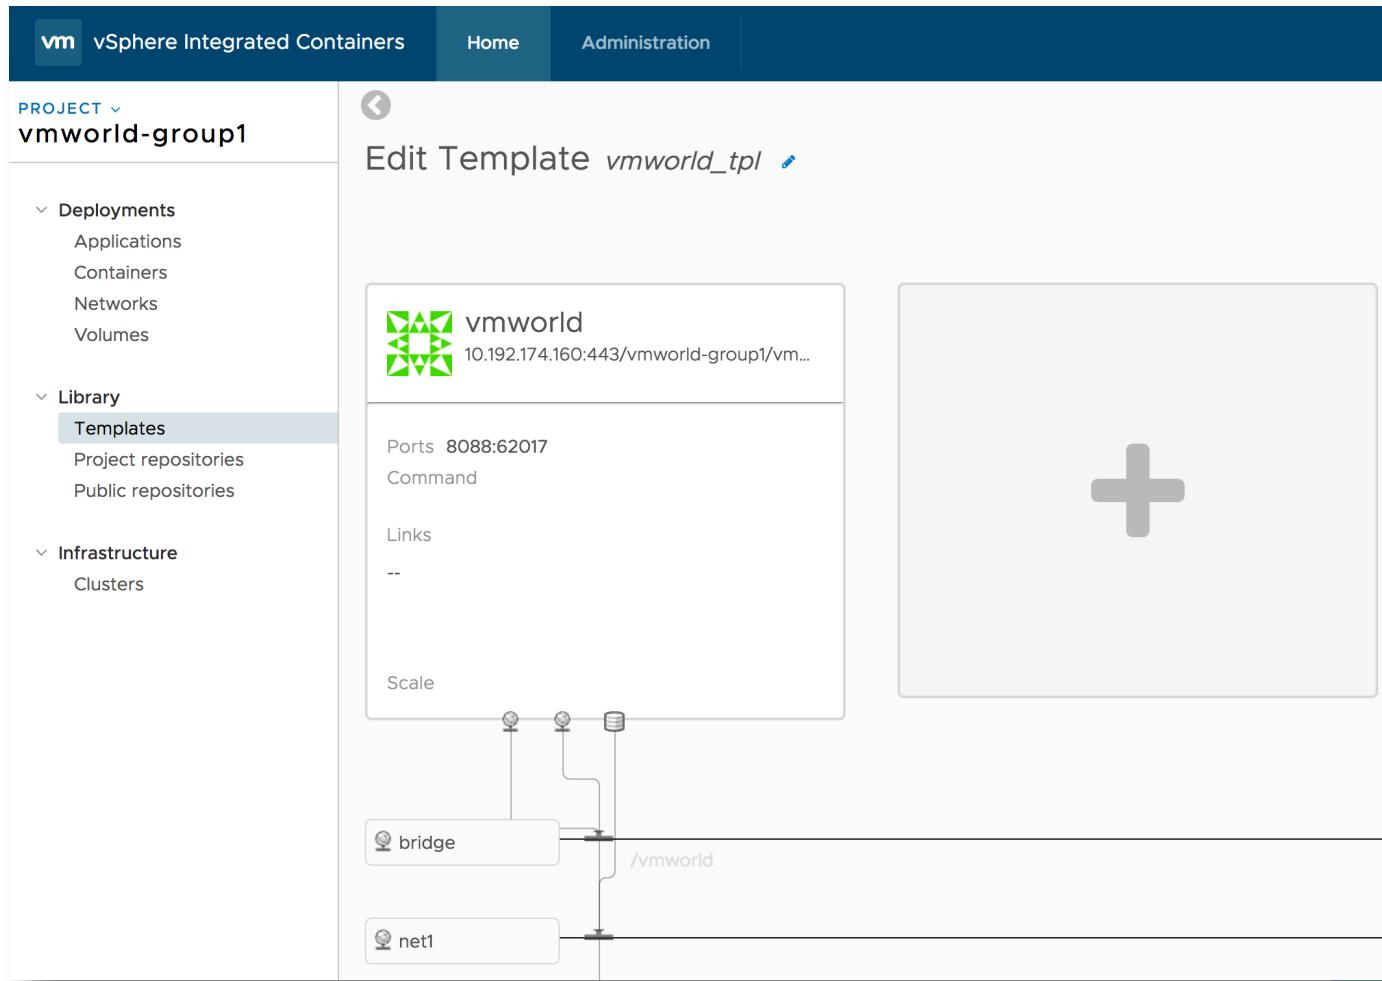
- Basic** tab selected.
- Image \***: 10.192.174.160:443/vmworld-group1/vmworld
- Name \***: myContainer
- Command**: Example: /startup.sh
- Links**: Service: [redacted], Alias: db
- Buttons**: PROVISION (blue), SAVE AS TEMPLATE (white)

# Application Template



ADMIRAL™

- Define your own application structure
  - Containers
  - Networks
  - Volumes
- Provision multiple containers at one time
- Compatible with **docker-compose** yaml file



# **Practice on VIC engine, Harbor & Admiral**

# Instructions

- Wifi: **vmworld\_hackathon\_training1 / Hackathon2017**
  - Download the practice manual book **VIC-lab-manualV4.pdf** from :  
<https://192.168.101.10:9443/files>
  - Practice group:
    - Project: vmworld\_group**1** – vmworld\_group**6**
    - User: vmworld\_u**1** – vmworld\_u**6**
  - Two practices included:
    - Basic Practice: Deploy container with specified image
    - Advanced Practice (optional): Deploy an application with template
- vmworld®**

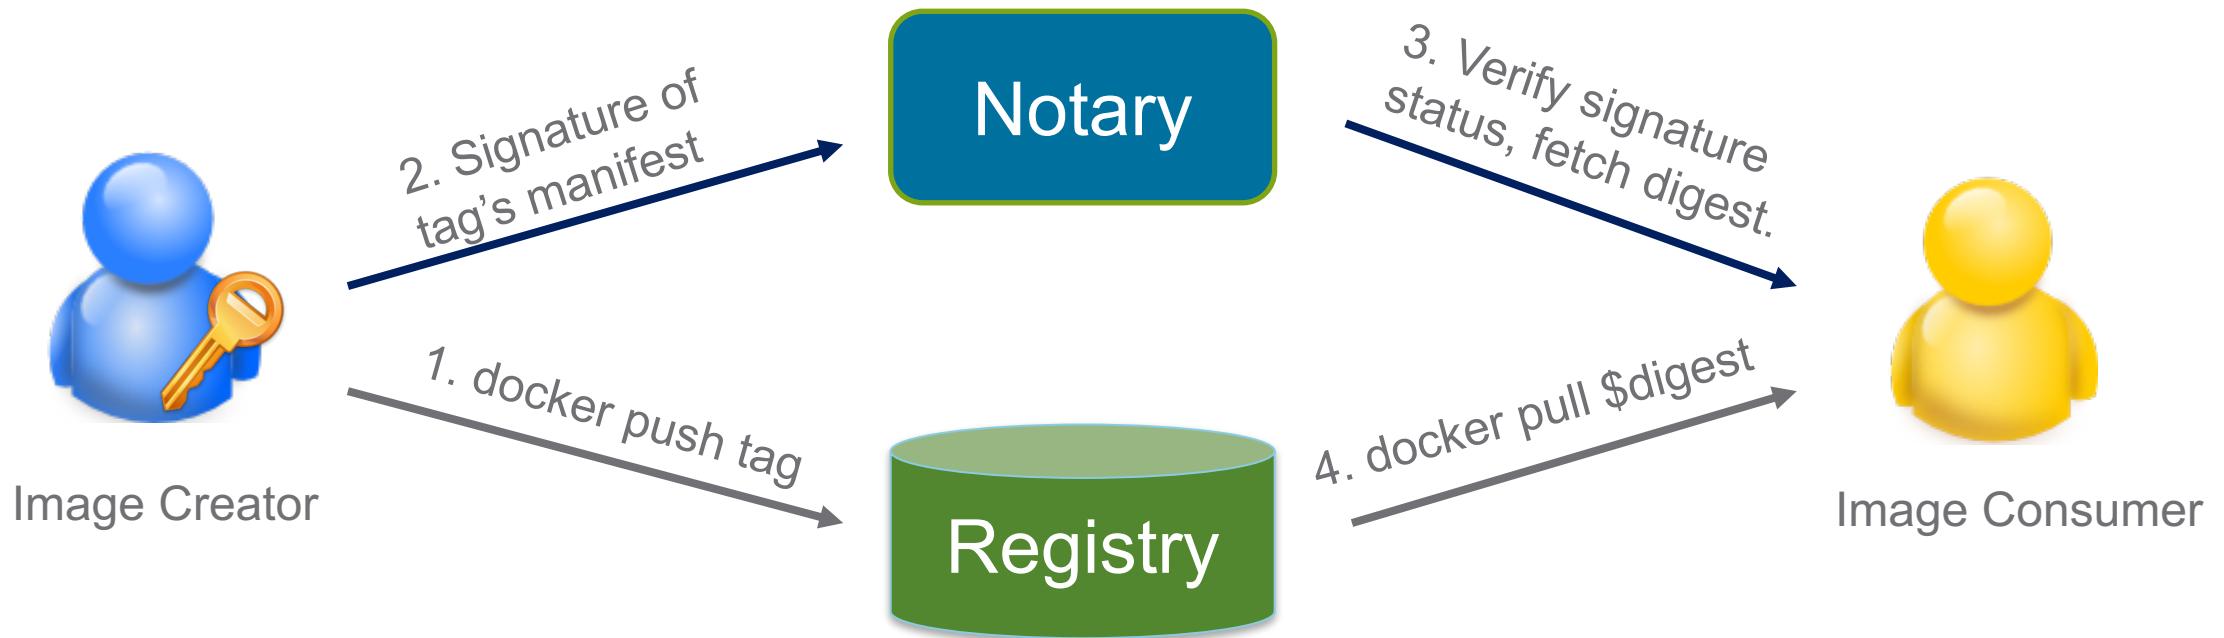
# Thank You

vmworld<sup>®</sup>  
2017

<https://github.com/vmware/harbor>

vmware<sup>®</sup>

# Content trust for image provenance



# Vulnerability Scanning

- **Static analysis** of vulnerability by inspecting filesystem of container image and indexing features in database.
- **Rescanning** is needed only and only if new detectors are added.
- Update vulnerability data regularly
  - Debian Security Bug Tracker
  - Ubuntu CVE Tracker
  - Red Hat Security Data
  - Oracle Linux Security Data
  - Alpine SecDB